

Lecture 1 0912

JongCheolLee

2024-09-02

Data1

```
## Open the dataset linked to the book website
url.ad <- "https://www.statlearning.com/s/Advertising.csv"
Advertising <- read.csv(url.ad, h=T)
attach(Advertising)
```

- attach() R ,


```
head(Advertising)
```

```
##   X      TV radio newspaper sales
## 1 1 230.1  37.8      69.2   22.1
## 2 2  44.5  39.3      45.1   10.4
## 3 3  17.2  45.9      69.3    9.3
## 4 4 151.5  41.3      58.5   18.5
## 5 5 180.8  10.8      58.4   12.9
## 6 6   8.7  48.9      75.0    7.2
```

```
summary(Advertising)
```

```
##           X              TV              radio      newspaper
## Min.      : 1.00    Min.      : 0.70    Min.      : 0.000    Min.      : 0.30
## 1st Qu.: 50.75    1st Qu.: 74.38    1st Qu.: 9.975    1st Qu.: 12.75
## Median :100.50    Median :149.75    Median :22.900    Median : 25.75
## Mean      :100.50    Mean      :147.04    Mean      :23.264    Mean      : 30.55
## 3rd Qu.:150.25    3rd Qu.:218.82    3rd Qu.:36.525    3rd Qu.: 45.10
## Max.      :200.00    Max.      :296.40    Max.      :49.600    Max.      :114.00
##      sales
## Min.      : 1.60
## 1st Qu.:10.38
## Median :12.90
## Mean      :14.02
## 3rd Qu.:17.40
## Max.      :27.00
```

```
## Least square fit for simple linear regression
par(mfrow = c(1,3))
plot(sales~TV, col=2, xlab="TV", ylab="Sales")
abline(lm(sales~TV)$coef, lwd=3, col="darkblue")
plot(sales~radio, col=2, xlab="Radio", ylab="Sales")
abline(lm(sales~radio)$coef, lwd=3, col="darkblue")
plot(sales~newspaper, col=2, xlab="Newspaper", ylab="Sales")
abline(lm(sales~newspaper)$coef, lwd=3, col="darkblue")
```



Multiple LR

```
AD <- Advertising[ , -1] #
```

```
## Multiple linear regression
lm.fit <- lm(sales ~ ., AD)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = sales ~ ., data = AD)
##
## Residuals:
```

```
##      Min      1Q  Median      3Q      Max
## -8.8277 -0.8908  0.2418  1.1893  2.8292
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.938889   0.311908   9.422  <2e-16 ***
## TV           0.045765   0.001395  32.809  <2e-16 ***
## radio        0.188530   0.008611  21.893  <2e-16 ***
## newspaper   -0.001037   0.005871  -0.177    0.86
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.686 on 196 degrees of freedom
## Multiple R-squared:  0.8972, Adjusted R-squared:  0.8956
## F-statistic: 570.3 on 3 and 196 DF,  p-value: < 2.2e-16
```

```
names(lm.fit)
```

```
## [1] "coefficients" "residuals"      "effects"      "rank"
## [5] "fitted.values" "assign"          "qr"           "df.residual"
## [9] "xlevels"       "call"            "terms"        "model"
```

```
coef(lm.fit)
```

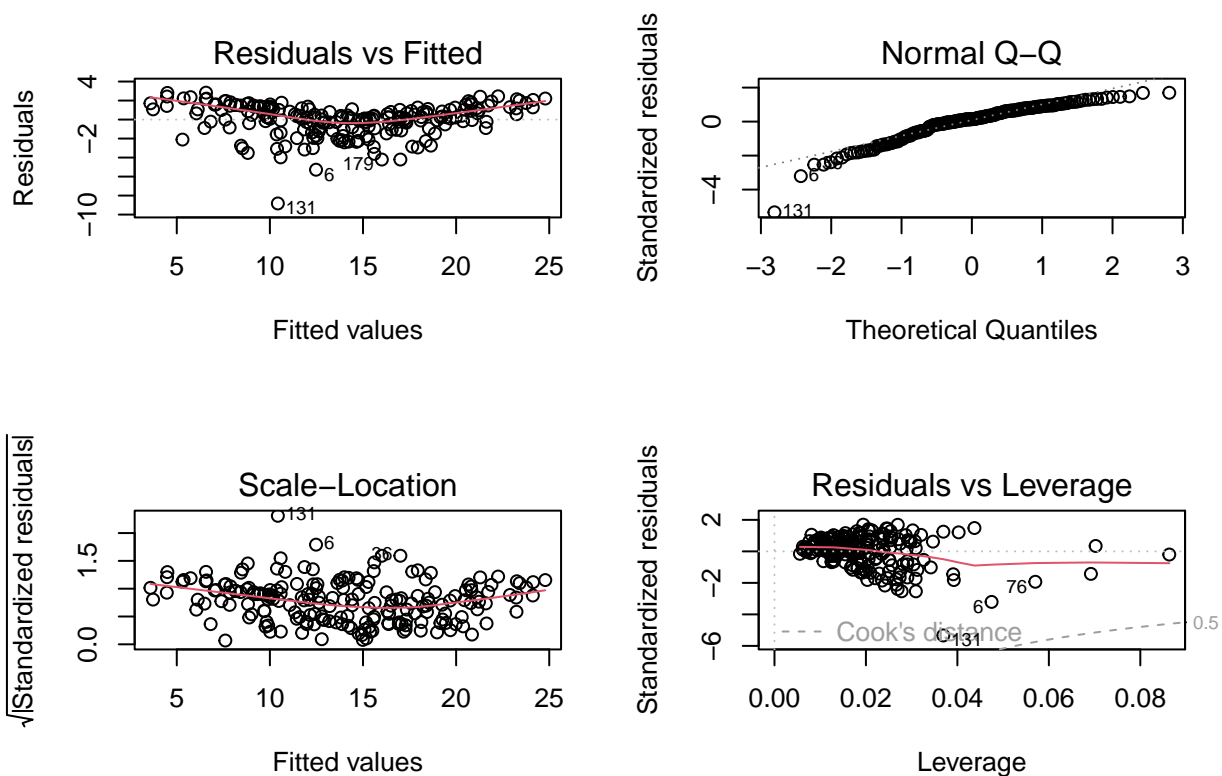
```
## (Intercept)      TV      radio  newspaper
## 2.938889369 0.045764645 0.188530017 -0.001037493
```

```
confint(lm.fit)
```

```
##              2.5 %      97.5 %
## (Intercept) 2.32376228 3.55401646
## TV          0.04301371 0.04851558
## radio       0.17154745 0.20551259
## newspaper   -0.01261595 0.01054097
```

- Newspaper p Newspaper

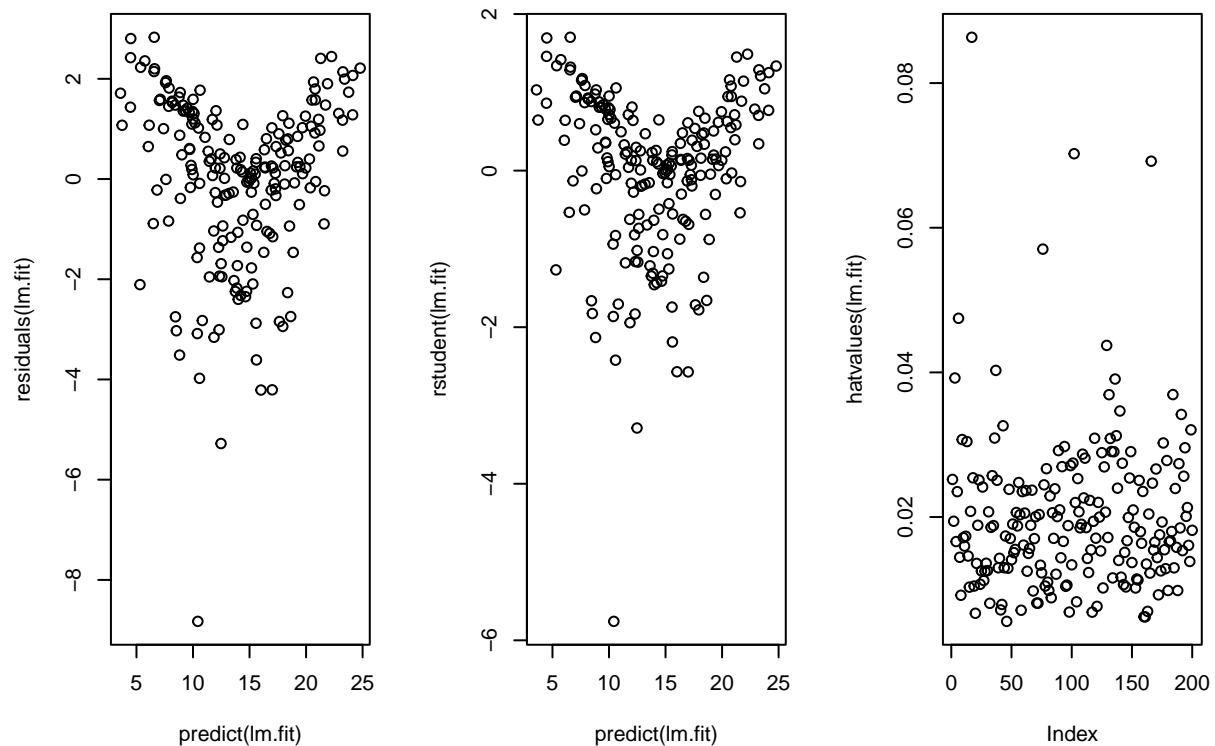
```
par(mfrow=c(2,2))
plot(lm.fit)
```



```
#dev.off()
```

- Residual Plot:
- Q-Q plot:
- Standardized Residuals:
- Studentized Residuals: $(\hat{e}_i) / \sqrt{1 - h_{ii}}$. Leverage
- Leverage (= Hat value):
 - Hat matrix
 - Hat value
 - $\hat{Y}_{\text{hat}} = H Y$, i hat value i
 - (\hat{Y}_{hat})
 - matrix
 - Leverage
 - Leverage

```
par(mfrow=c(1,3))
plot(predict(lm.fit), residuals(lm.fit))
plot(predict(lm.fit), rstudent(lm.fit))
plot(hatvalues(lm.fit))
```



```
which.max(hatvalues(lm.fit))
```

```
## 17
## 17
```

```
• hat          .   hat      ,          .
```

Data2

```
url.in <- "https://www.statlearning.com/s/Income1.csv"
Income <- read.csv(url.in, h=T)
attach(Income)
```

```
## The following object is masked _by_ .GlobalEnv:
##
##      Income

## The following object is masked from Advertising:
##
##      X
```

```
head(Income)
```

```
##   X Education   Income
## 1 1  10.00000 26.65884
## 2 2  10.40134 27.30644
## 3 3  10.84281 22.13241
## 4 4  11.24415 21.16984
## 5 5  11.64548 15.19263
## 6 6  12.08696 26.39895
```



```
summary(Income)
```

```
##           X           Education           Income
## Min.      : 1.00   Min.      :10.00   Min.      :15.19
## 1st Qu.:  8.25   1st Qu.:12.99   1st Qu.:29.08
## Median :15.50   Median :16.00   Median :49.87
## Mean    :15.50   Mean    :16.00   Mean    :50.15
## 3rd Qu.:22.75   3rd Qu.:19.01   3rd Qu.:71.14
## Max.    :30.00   Max.    :22.00   Max.    :80.26
```

Polynomial regression

```
## Polynomial regression fit
par(mfrow = c(1,2))
plot(Income~Education, col=2, pch=19, xlab="Years of Education",
ylab="Income", data=Income)
g <- lm(Income ~ poly(Education, 3), data=Income)

plot(Income~Education, col=2, pch=19, xlab="Years of Education",
ylab="Income", data=Income)
lines(Income$Education, g$fit, col="darkblue", lwd=4,
ylab="Income", xlab="Years of Education")
```



```
y <- Income$Income
mean((predict(g) - y)^2) # mean(residuals(g)^2) .
```

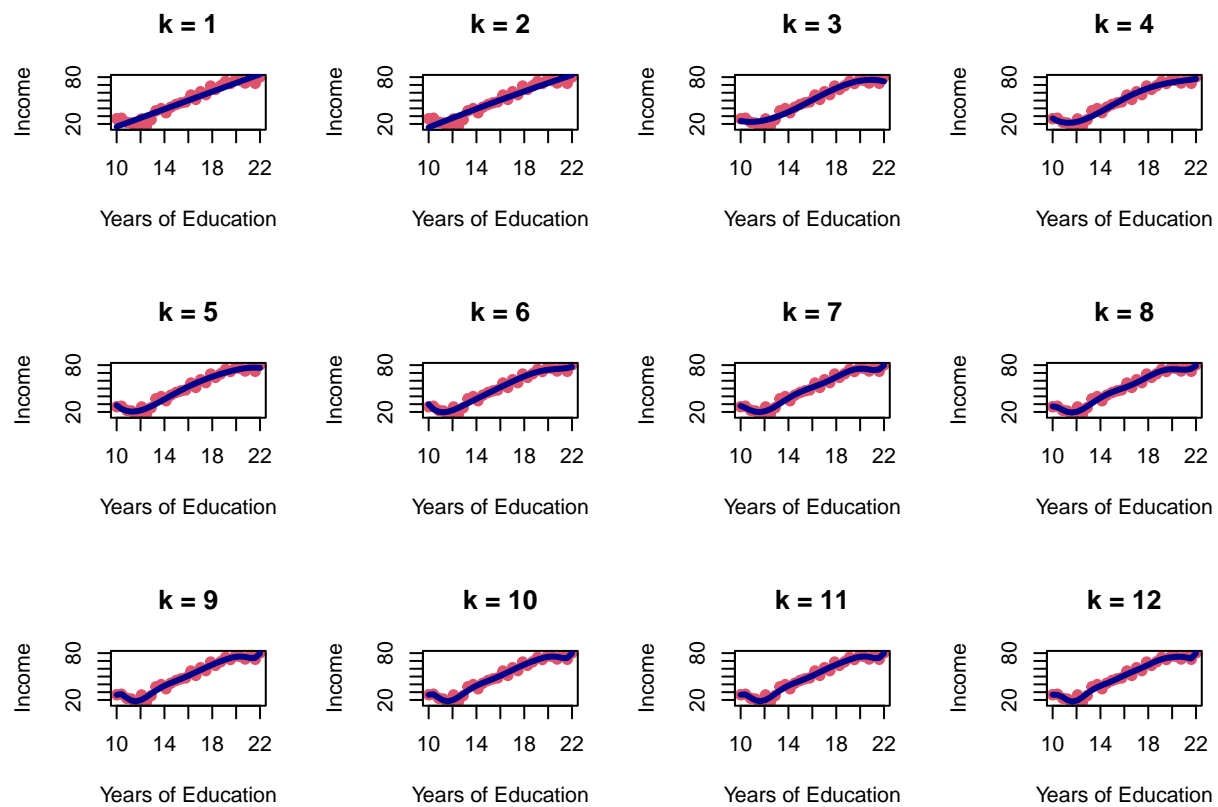
```
## [1] 15.10126
```

```
sum(predict(g) - y)
```


```
## [1] -1.918465e-13
```

Polynomial regression from deg 1 to 12

```
dist <- NULL
par(mfrow=c(3,4))
for (k in 1:12) {
  g <- lm(Income ~ poly(Education, k), data=Income)
  dist[k] <- mean(residuals(g)^2)
  plot(Income~Education, col=2, pch=19,
        xlab="Years of Education", ylab="Income",
        data=Income, main=paste("k =", k))
  lines(Income$Education,g$fit,col="darkblue",lwd=3,
        ylab="Income", xlab="Years of Education")
}
```



```
#x11()
plot(dist, type="b", xlab="Degree of Polynomial",
      ylab="Mean squared distance")
```

```
dist # MSE
```

```
## [1] 29.828816 29.590053 15.101265 12.675769 12.221552 11.680070 10.039115
## [8] 9.825701 9.276918 9.265054 9.254865 9.064854
```

- 12 MSE . ? => Overfitting .

Training measurement vs Test measurement

```
set.seed(12345)
## Simulate x and y based on a known function
fun1 <- function(x) -(x-100)*(x-30)*(x+15)/13^4+6 # True underlying model
x <- runif(50,0,100)
y <- fun1(x) + rnorm(50) #

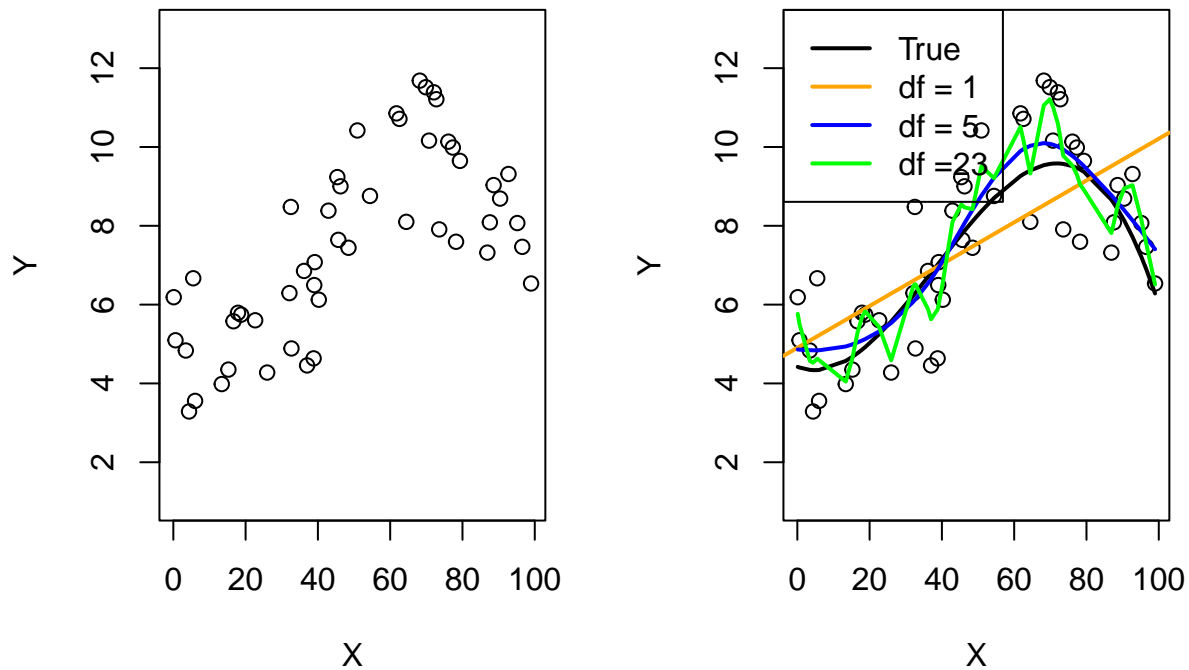
## Plot linear regression and splines (Prediction models)
par(mfrow=c(1,2))
plot(x, y, xlab="X", ylab="Y", ylim=c(1,13))
plot(x, y, xlab="X", ylab="Y", ylim=c(1,13))
lines(sort(x), fun1(sort(x)), col=1, lwd=2)
abline(lm(y~x)$coef, col="orange", lwd=2)

lines(smooth.spline(x,y, df=5), col="blue", lwd=2) # smoothing spline
```

```

lines(smooth.spline(x,y, df=23), col="green", lwd=2)
legend("topleft", lty=1, col=c(1, "orange", "blue", "green"),
legend=c("True", "df = 1", "df = 5", "df =23"),lwd=2)

```



```

set.seed(45678)
## Simulate training and test data (x, y)
tran.x <- runif(50,0,100)
test.x <- runif(50,0,100)
tran.y <- fun1(tran.x) + rnorm(50)
test.y <- fun1(test.x) + rnorm(50)

## Compute MSE along with different df
df <- 2:40
MSE <- matrix(0, length(df), 2)

for (i in 1:length(df)) {
  tran.fit <- smooth.spline(tran.x, tran.y, df=df[i])
  MSE[i,1] <- mean((tran.y - predict(tran.fit, tran.x)$y)^2) # training set
  MSE[i,2] <- mean((test.y - predict(tran.fit, test.x)$y)^2) # test set
}

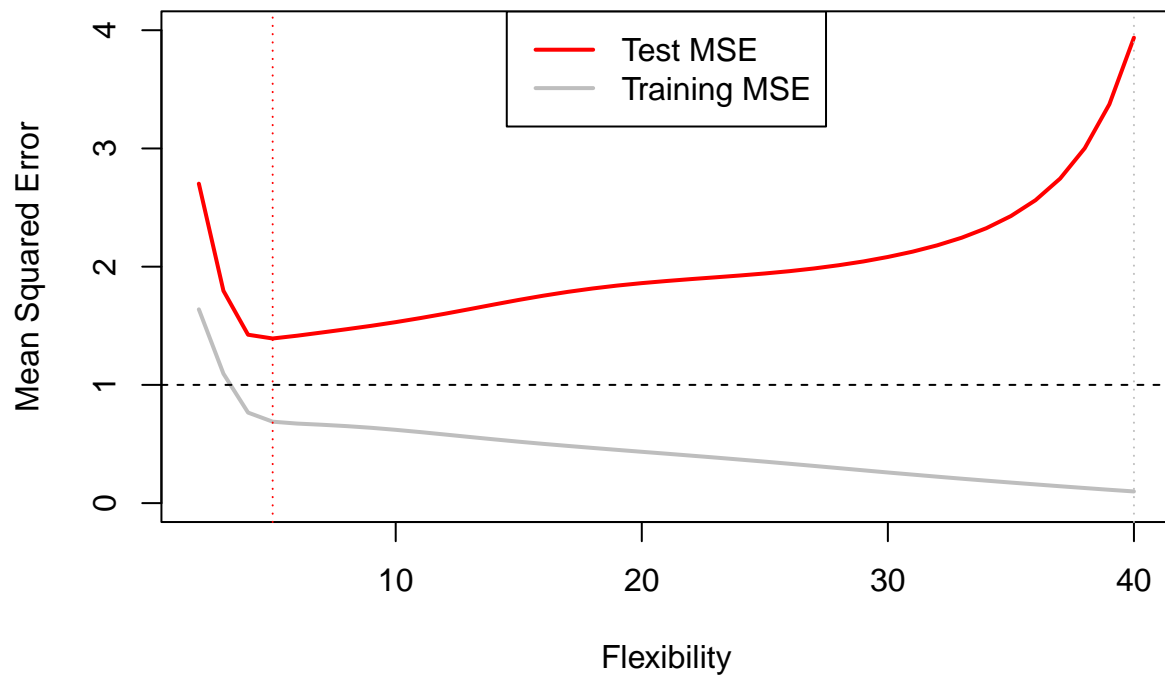
## Plot both test and training errors
matplot(df, MSE, type="l", col=c("gray", "red"),
xlab="Flexibility", ylab="Mean Squared Error",
lwd=2, lty=1, ylim=c(0,4))

```

```

abline(h=1, lty=2)
legend("top", lty=1, col=c("red", "gray"),lwd=2,
legend=c("Test MSE", "Training MSE"))
abline(v=df[which.min(MSE[,1])], lty=3, col="gray")
abline(v=df[which.min(MSE[,2])], lty=3, col="red")


```



```

matplot(df, MSE, type="l", col=c("gray", "red"),
xlab="Flexibility", ylab="Mean Squared Error",
lwd=2, lty=1, ylim=c(0.5,2), xlim=c(1,10))
abline(h=1, lty=2)
abline(v=df[which.min(MSE[,2])], lty=3, col="red")



```



- Training MSE
- (Unseen) Test MSE, training set

```
set.seed(12345)
## Simulate x and y based on a known function
fun2 <- function(x) x/10 +2 # true underlying model: Linear
x <- runif(50,0,100)
y <- fun2(x) + rnorm(50)

## Plot linear regression and splines
par(mfrow=c(1,2))
plot(x, y, xlab="X", ylab="Y", ylim=c(1,13))
plot(x, y, xlab="X", ylab="Y", ylim=c(1,13))
lines(sort(x), fun2(sort(x)), col=1, lwd=2)
abline(lm(y~x)$coef, col="orange", lwd=2)
lines(smooth.spline(x,y, df=5), col="blue", lwd=2)
lines(smooth.spline(x,y, df=23), col="green", lwd=2)
legend("topleft", lty=1, col=c(1, "orange", "blue", "green"),
legend=c("True", "df = 1", "df = 5", "df =23"),lwd=2)
```

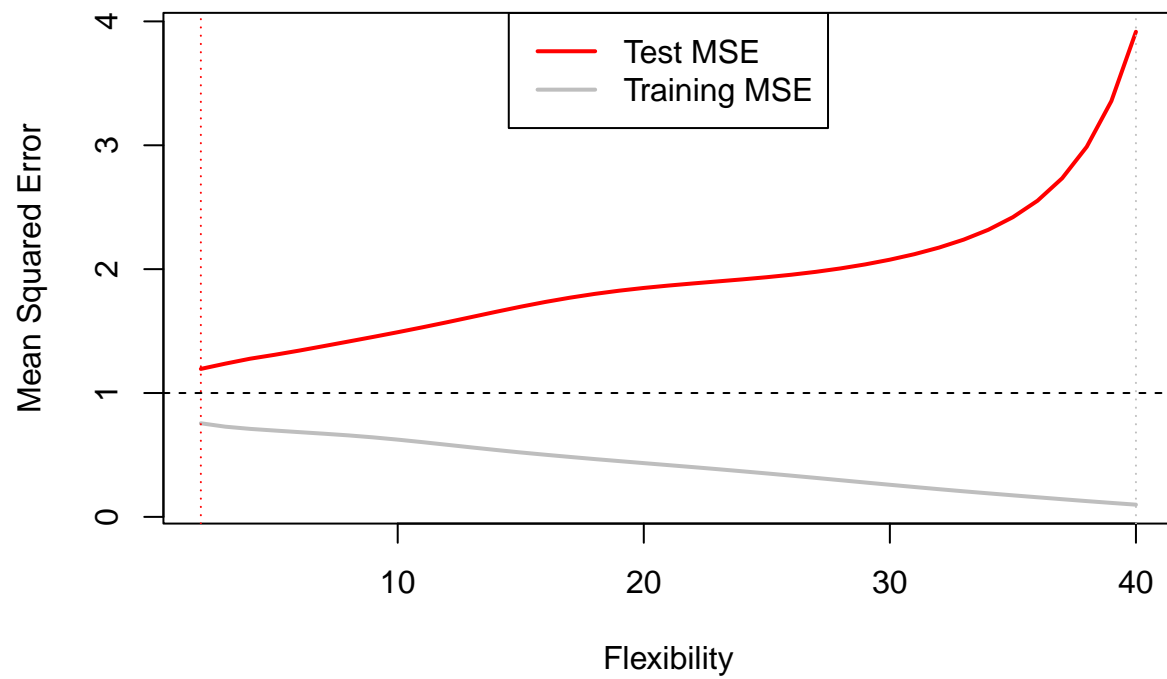


```
set.seed(45678)


## Simulate training and test data (x, y)
tran.x <- runif(50,0,100)
test.x <- runif(50,0,100)
tran.y <- fun2(tran.x) + rnorm(50)
test.y <- fun2(test.x) + rnorm(50)

## Compute MSE along with different df
df <- 2:40
MSE <- matrix(0, length(df), 2)
for (i in 1:length(df)) {
  tran.fit <- smooth.spline(tran.x, tran.y, df=df[i]) # training data x,y
  MSE[i,1] <- mean((tran.y - predict(tran.fit, tran.x)$y)^2)
  MSE[i,2] <- mean((test.y - predict(tran.fit, test.x)$y)^2)
}

## Plot both test and training errors
matplot(df, MSE, type="l", col=c("gray", "red"),
        xlab="Flexibility", ylab="Mean Squared Error",
        lwd=2, lty=1)
abline(h=1, lty=2)
legend("top", lty=1, col=c("red", "gray"),lwd=2,
       legend=c("Test MSE", "Training MSE"))
abline(v=df[which.min(MSE[,1])], lty=3, col="gray")
abline(v=df[which.min(MSE[,2])], lty=3, col="red")
```





```
matplot(df, MSE, type="l", col=c("gray", "red"),
xlab="Flexibility", ylab="Mean Squared Error",
lwd=2, lty=1, ylim=c(0,1.5), xlim=c(1,10))
abline(h=1, lty=2)
abline(v=df[which.min(MSE[,2])], lty=3, col="red")
```



```
set.seed(12345)
## Simulate x and y based on a known function
fun3 <- function(x) -(x-80)*(x-45)*(x-25)/15^3+10
x <- runif(50,0,100)
y <- fun3(x) + rnorm(50)


## Plot linear regression and splines
par(mfrow=c(1,2))
plot(x, y, xlab="X", ylab="Y")
plot(x, y, xlab="X", ylab="Y")
lines(sort(x), fun3(sort(x)), col=1, lwd=2)
abline(lm(y~x)$coef, col="orange", lwd=2)
lines(smooth.spline(x,y, df=5), col="blue", lwd=2)
lines(smooth.spline(x,y, df=23), col="green", lwd=2)
legend("topright", lty=1, col=c(1, "orange", "blue", "green"),
legend=c("True", "df = 1", "df = 5", "df =23"),lwd=2)
```




```
set.seed(45678)
## Simulate training and test data (x, y)
tran.x <- runif(50,0,100)
test.x <- runif(50,0,100)
tran.y <- fun3(tran.x) + rnorm(50)
test.y <- fun3(test.x) + rnorm(50)

## Compute MSE along with different df
df <- 2:40
MSE <- matrix(0, length(df), 2)
for (i in 1:length(df)) {
  tran.fit <- smooth.spline(tran.x, tran.y, df=df[i])
  MSE[i,1] <- mean((tran.y - predict(tran.fit, tran.x)$y)^2)
  MSE[i,2] <- mean((test.y - predict(tran.fit, test.x)$y)^2)
}
```

```
## Plot both test and training errors
matplot(df, MSE, type="l", col=c("gray", "red"),
        xlab="Flexibility", ylab="Mean Squared Error",
        lwd=2, lty=1)
abline(h=1, lty=2)
legend("top", lty=1, col=c("red", "gray"), lwd=2,
       legend=c("Test MSE", "Training MSE"))
abline(v=df[which.min(MSE[,1])], lty=3, col="gray")
abline(v=df[which.min(MSE[,2])], lty=3, col="red")
```

```
matplot(df, MSE, type="l", col=c("gray", "red"),
xlab="Flexibility", ylab="Mean Squared Error",
lwd=2, lty=1, ylim=c(0,3), xlim=c(4,20))
abline(h=1, lty=2)
abline(v=df[which.min(MSE[,2])], lty=3, col="red")
```



Validation Set Approach

```
library(ISLR)
```

```
## Warning: 'ISLR' R 4.2.3
```

```
data(Auto)
str(Auto)
```

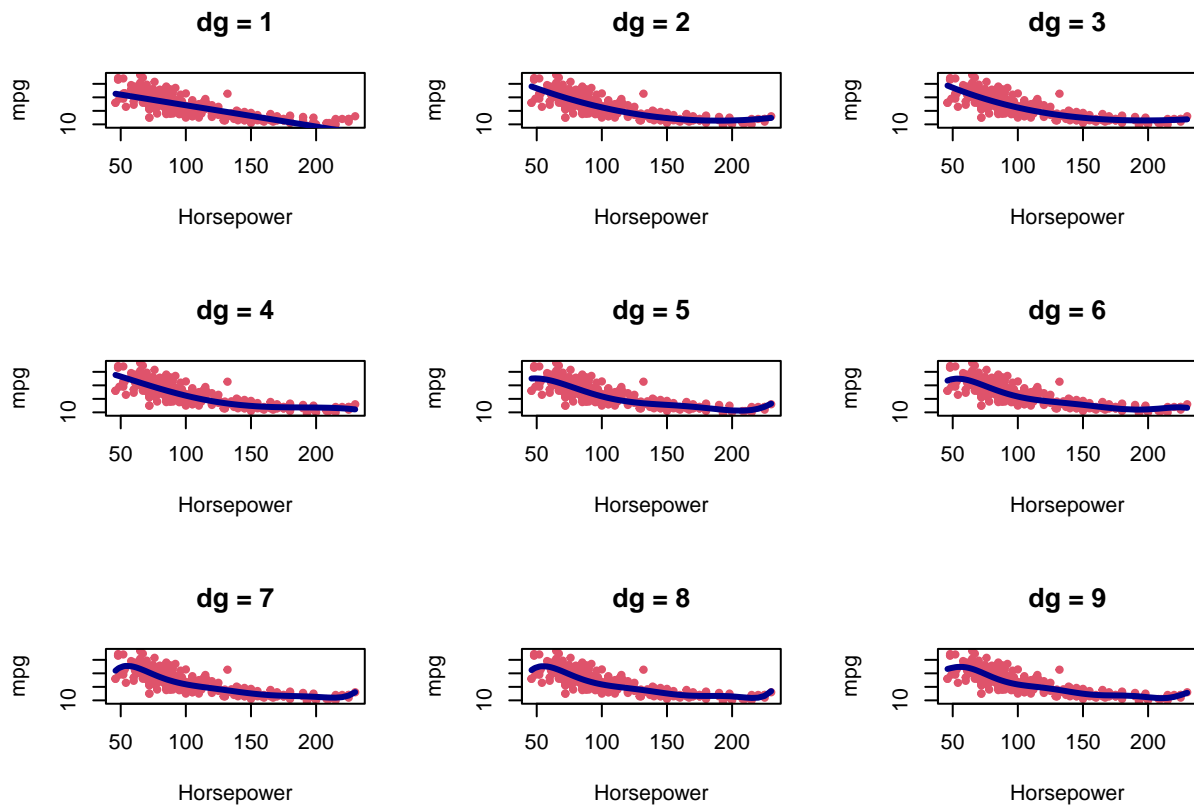
```
## 'data.frame': 392 obs. of 9 variables:
## $ mpg : num 18 15 18 16 17 15 14 14 15 ...
## $ cylinders : num 8 8 8 8 8 8 8 8 8 ...
## $ displacement: num 307 350 318 304 302 429 454 440 455 390 ...
## $ horsepower : num 130 165 150 150 140 198 220 215 225 190 ...
## $ weight : num 3504 3693 3436 3433 3449 ...
## $ acceleration: num 12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
## $ year : num 70 70 70 70 70 70 70 70 70 ...
## $ origin : num 1 1 1 1 1 1 1 1 1 ...
## $ name : Factor w/ 304 levels "amc ambassador brougham",...: 49 36 231 14 161 141 54 223 241 1
```

```
summary(Auto)
```

```
##      mpg      cylinders    displacement    horsepower      weight
## Min.   : 9.00    Min.   :3.000    Min.   : 68.0    Min.   : 46.0    Min.   :1613
## 1st Qu.:17.00    1st Qu.:4.000    1st Qu.:105.0    1st Qu.: 75.0    1st Qu.:2225
## Median :22.75    Median :4.000    Median :151.0    Median : 93.5    Median :2804
## Mean   :23.45    Mean   :5.472    Mean   :194.4    Mean   :104.5    Mean   :2978
## 3rd Qu.:29.00    3rd Qu.:8.000    3rd Qu.:275.8    3rd Qu.:126.0    3rd Qu.:3615
## Max.   :46.60    Max.   :8.000    Max.   :455.0    Max.   :230.0    Max.   :5140
##
##      acceleration      year      origin      name
## Min.   : 8.00    Min.   :70.00    Min.   :1.000    amc matador      : 5
## 1st Qu.:13.78    1st Qu.:73.00    1st Qu.:1.000    ford pinto       : 5
## Median :15.50    Median :76.00    Median :1.000    toyota corolla   : 5
## Mean   :15.54    Mean   :75.98    Mean   :1.577    amc gremlin      : 4
## 3rd Qu.:17.02    3rd Qu.:79.00    3rd Qu.:2.000    amc hornet       : 4
## Max.   :24.80    Max.   :82.00    Max.   :3.000    chevrolet chevette: 4
##                                     (Other)      :365
```

```
mpg <- Auto$mpg
horsepower <- Auto$horsepower
dg <- 1:9
u <- order(horsepower)

par(mfrow=c(3,3))
for (k in 1:length(dg)) {
  g <- lm(mpg ~ poly(horsepower, dg[k]))
  plot(mpg~horsepower, col=2, pch=20, xlab="Horsepower",
       ylab="mpg", main=paste("dg =", dg[k]))
  lines(horsepower[u], g$fit[u], col="darkblue", lwd=3)
}
```



```
set.seed(1)
n <- nrow(Auto)

## training set
tran <- sample(n, n/2) #
MSE <- NULL

for (k in 1:length(dg)) {
  g <- lm(mpg ~ poly(horsepower, dg[k]), subset=tran) # training set
  MSE[k] <- mean((mpg - predict(g, Auto))[-tran]^2) # MSE
}

par(mfrow=c(1,3))
plot(dg, MSE, type="b", col=2, xlab="Degree of Polynomial",
     ylab="Mean Squared Error", ylim=c(15,30), lwd=2, pch=19)
abline(v=which.min(MSE), lty=2)

K <- 10
MSE <- matrix(0, length(dg), K)

for (i in 1:K) {
  tran <- sample(392, 196) # K training set (K )
  for (k in 1:length(dg)) { #
    g <- lm(mpg ~ poly(horsepower, dg[k]), subset=tran)
    MSE[k, i] <- mean((mpg - predict(g, Auto))[-tran]^2)
  }
}
```

```


}

matplot(dg, MSE, type="l", xlab="Degree of Polynomial", lty=1,
        ylab="Mean Squared Error", col=1:10, ylim=c(15,30))

avg <- apply(MSE, 1, mean) # K

plot(dg, avg, type="b", col=2, xlab="Degree of Polynomial",
     ylab="Mean Squared Error", ylim=c(15,30), lwd=2, pch=19)
abline(v=which.min(avg), lty=2)

```



- training set .

Leave one out cross validation

```

n <- nrow(Auto)
dg <- 1:9
MSE <- matrix(0, n, length(dg))

for (i in 1:n) { # validation
  for (k in 1:length(dg)) {
    g <- lm(mpg ~ poly(horsepower, k), subset=(1:n)[-i]) # i
    MSE[i, k] <- mean((mpg - predict(g, Auto))[i]^2) # i mse
  }
}

```

```

    }
  }

aMSE <- apply(MSE, 2, mean)

par(mfrow=c(1, 2))
plot(dg, aMSE, type="b", col="darkblue",
     xlab="Degree of Polynomial", ylab="Mean Squared Error",
     ylim=c(18,25), lwd=2, pch=19)
abline(v=which.min(aMSE), lty=2)

ncv <- NULL

for (k in 1:length(dg)) {
  g <- lm(mpg ~ poly(horsepower, k))
  ncv[k] <- mean((g$res/(1-influence(g)$hat))^2) # influence()      leverage
  # for loop n*k
}
lines(dg, ncv, col=2, lty=2, lwd=2)

K <- 10 ## 10-fold cross validation.      obs      10      .
MSE <- matrix(0, n, length(dg))

set.seed(54321)
u <- sample(rep(seq(K), length=n))
table(u)

```

```

## u
##  1  2  3  4  5  6  7  8  9 10
## 40 40 39 39 39 39 39 39 39 39

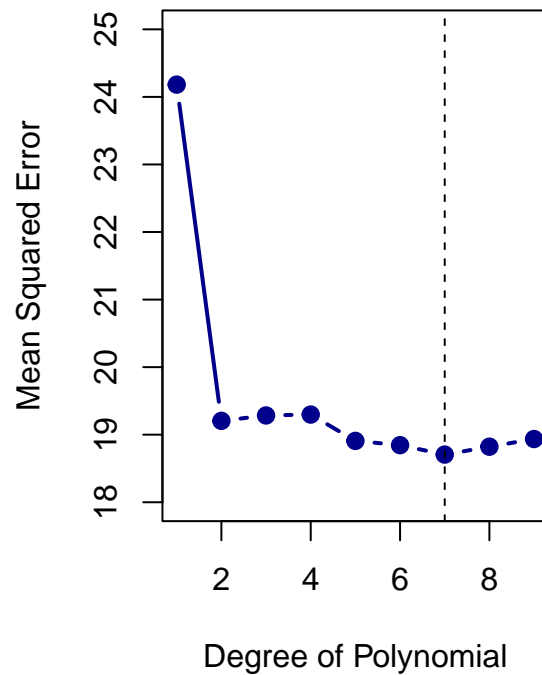

```

```

for (k in 1:K) {
  tran <- which(u!=k)
  test <- which(u==k)
  for (i in 1:length(dg)) {
    g <- lm(mpg ~ poly(horsepower, i), subset=tran)
    MSE[test, i] <- (mpg - predict(g, Auto))[test]^2
  }
}
CVE <- apply(MSE, 2, mean)

plot(dg, CVE, type="b", col="darkblue",
     xlab="Degree of Polynomial", ylab="Mean Squared Error",
     ylim=c(18,25), lwd=2, pch=19)
abline(v=which.min(CVE), lty=2)

```



K-fold CV

```


N <- 9 ## Number of K-fold CV replications
KCV <- matrix(0, length(dg), N)

set.seed(1234)
for (j in 1:N) {
  MSE <- matrix(0, n, length(dg))
  u <- sample(rep(seq(K), length=n))
  for (k in 1:K) {
    tran <- which(u!=k)
    test <- which(u==k)
    for (i in 1:length(dg)) {
      g <- lm(mpg ~ poly(horsepower, i), subset=tran)
      MSE[test, i] <- (mpg - predict(g, Auto))[test]^2
    }
  }
  KCV[,j] <- apply(MSE, 2, mean)
}

matplot(dg, KCV, type="l", xlab="Degree of Polynomial", lty=1,
ylab="MSE", ylim=c(18,25), main="10-fold CV")

```

10-fold CV



```
apply(KCV, 2, which.min)
```

```
## [1] 7 7 7 7 7 7 7 7 7
```

- K . => K-fold CV .

boot package : CV

```
library(boot)
set.seed(101010)

## Leave-one-out CV
MSE <- NULL
for (i in 1:length(dg)) {
  glm.fit <- glm(mpg ~ poly(horsepower ,i))
  MSE[i] <- cv.glm(Auto, glm.fit)$delta[1]
}
plot(dg, MSE, type="b", col="darkblue", ylim=c(15,29),
      xlab="Degree of Polynomial", ylab="MSE", lwd=2, pch=19)


## K-fold cross validation
K <- 10
KCV <- NULL
```



```

for (i in 1:length(dg)) {
  glm.fit <- glm(mpg ~ poly(horsepower ,i))
  KCV[i] <- cv.glm(Auto, glm.fit, K=K)$delta[1]
}
lines(dg, KCV, col=2, lwd=2, type="b", pch=19, lty=2)

```



Cubic Model with K-fold Cross Validation

- training set df MSE , CV .

```

set.seed(45678)
x <- runif(50,0,100)
y <- fun1(x) + rnorm(50) # cubic function + noise

K <- 5
df <- 2:40 # 39

MSE <- matrix(0, length(x), length(df))
u <- sample(rep(seq(K), length=length(x))) # 1:K

for (k in 1:K) {
  tr <- which(u!=k)
  te <- which(u==k)
  for (j in 1:length(df)) {

```

```

    fit <- smooth.spline(x[tr], y[tr], df=df[j])
    MSE[te, j] <- y[te] - predict(fit, x[te])$y
  }
}

CVE <- apply(MSE^2, 2, mean)
data.frame(DF=df, CVE=CVE)

```

```

##      DF      CVE
## 1     2  1.826783
## 2     3  1.149435
## 3     4  1.050967
## 4     5  1.082088
## 5     6  1.145289
## 6     7  1.225121
## 7     8  1.309276
## 8     9  1.388198
## 9    10  1.456269
## 10   11  1.513155
## 11   12  1.561547
## 12   13  1.605483
## 13   14  1.649503
## 14   15  1.697560
## 15   16  1.752622
## 16   17  1.816434
## 17   18  1.889711
## 18   19  1.973275
## 19   20  2.066752
## 20   21  2.169945
## 21   22  2.283464
## 22   23  2.408992
## 23   24  2.551177
## 24   25  2.717405
## 25   26  2.918233
## 26   27  3.171596
## 27   28  3.493487
## 28   29  3.901390
## 29   30  4.410251
## 30   31  5.014618
## 31   32  5.684375
## 32   33  6.370503
## 33   34  7.015428
## 34   35  7.603206
## 35   36  8.240476
## 36   37  9.688811
## 37   38 15.674305
## 38   39 40.679175
## 39   40 113.227168

```

Linear Model with K-fold Cross Validation

```
set.seed(45678)
x <- runif(50,0,100)
y <- fun2(x) + rnorm(50)

K <- 5
df <- 2:40

MSE <- matrix(0, length(x), length(df))
u <- sample(rep(seq(K), length=length(x)))

for (k in 1:K) {
  tr <- which(u!=k)
  te <- which(u==k)
  for (j in 1:length(df)) {
    fit <- smooth.spline(x[tr], y[tr], df=df[j])
    MSE[te, j] <- y[te] - predict(fit, x[te])$y
  }
}

CVE <- apply(MSE^2, 2, mean)
data.frame(DF=df, CVE=CVE)
```

##	DF	CVE
## 1	2	1.037058
## 2	3	1.047150
## 3	4	1.061881
## 4	5	1.109273
## 5	6	1.175957
## 6	7	1.253769
## 7	8	1.334445
## 8	9	1.410614
## 9	10	1.476878
## 10	11	1.532417
## 11	12	1.579600
## 12	13	1.622357
## 13	14	1.665218
## 14	15	1.712177
## 15	16	1.766228
## 16	17	1.829125
## 17	18	1.901586
## 18	19	1.984428
## 19	20	2.077256
## 20	21	2.179865
## 21	22	2.292861
## 22	23	2.417904
## 23	24	2.559636
## 24	25	2.725429
## 25	26	2.925834
## 26	27	3.178782
## 27	28	3.500271
## 28	29	3.907779

```
## 29 30 4.416252
## 30 31 5.020225
## 31 32 5.689573
## 32 33 6.375282
## 33 34 7.019800
## 34 35 7.607257
## 35 36 8.244439
## 36 37 9.693230
## 37 38 15.680525
## 38 39 40.689656
## 39 40 113.244389
```

Nonlinear Model with K-fold Cross Validation

```
set.seed(45678)
x <- runif(50,0,100)
y <- fun3(x) + rnorm(50)

K <- 5
df <- 2:40

MSE <- matrix(0, length(x), length(df))
u <- sample(rep(seq(K), length=length(x)))

for (k in 1:K) {
  tr <- which(u!=k)
  te <- which(u==k)
  for (j in 1:length(df)) {
    fit <- smooth.spline(x[tr], y[tr], df=df[j])
    MSE[te, j] <- y[te] - predict(fit, x[te])$y
  }
}

CVE <- apply(MSE^2, 2, mean)
data.frame(DF=df, CVE=CVE)
```

```
##      DF      CVE
## 1    2 34.722550
## 2    3 21.947697
## 3    4  6.961509
## 4    5  2.682327
## 5    6  1.751877
## 6    7  1.491226
## 7    8  1.434082
## 8    9  1.443439
## 9   10  1.472216
## 10  11  1.505855
## 11  12  1.540830
## 12  13  1.577728
## 13  14  1.618553
## 14  15  1.665584
## 15  16  1.720771
```

##	16	17	1.785246
##	17	18	1.859386
##	18	19	1.943828
##	19	20	2.038145
##	20	21	2.142112
##	21	22	2.256258
##	22	23	2.382320
##	23	24	2.524912
##	24	25	2.691465
##	25	26	2.892544
##	26	27	3.146058
##	27	28	3.468014
##	28	29	3.875961
##	29	30	4.384904
##	30	31	4.989569
##	31	32	5.660000
##	32	33	6.347225
##	33	34	6.993501
##	34	35	7.582228
##	35	36	8.218605
##	36	37	9.661143
##	37	38	15.628178
##	38	39	40.589758
##	39	40	113.069328