

Lecture02

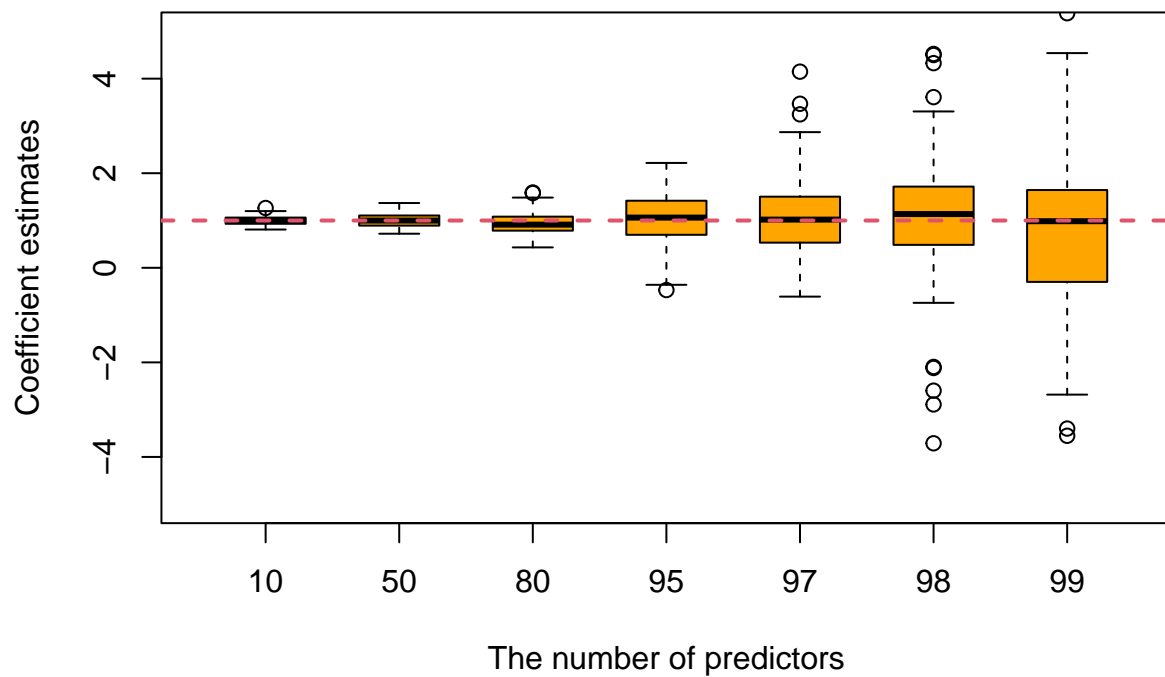
JongCheolLee

2024-09-23

```
set.seed(123)
n <- 100
pp <- c(10, 50, 80, 95, 97, 98, 99) #
B <- matrix(0, 100, length(pp))

for (i in 1:100) { # 100
  for (j in 1:length(pp)) {
    beta <- rep(0, pp[j])
    beta[1] <- 1 # Beta1          0. Bias = E(Beta1_hat) - 1
    x <- matrix(rnorm(n*pp[j]), n, pp[j])
    y <- x %*% beta + rnorm(n) # True Linear Model
    g <- lm(y~x) # Estimation
    B[i,j] <- g$coef[2] #          , Beta1_hat
  }
}

boxplot(B, col="orange", boxwex=0.6, ylab="Coefficient estimates",
names=pp, xlab="The number of predictors", ylim=c(-5,5))
abline(h=1, col=2, lty=2, lwd=2)
```



```
apply(B, 2, mean)
```

```
## [1] 1.0005277 1.0057437 0.9432175 1.0358624 1.0855411 1.0727110 0.9223116
```

```
• LSE unbiased      1
```

```
apply(B, 2, var)
```

```
## [1] 0.008622572 0.022278710 0.054041523 0.301848895 0.685693276
```

```
## [6] 5.941321509 38.018106760
```

```
•      B1      .
```

Best Subset Selection

```
library(ISLR)
```

```
## Warning: 'ISLR' R 4.2.3
```

```
names(Hitters)
```

```
## [1] "AtBat"      "Hits"       "HmRun"      "Runs"       "RBI"        "Walks"
## [7] "Years"      "CAAtBat"    "CHits"      "CHmRun"     "CRuns"      "CRBI"
## [13] "CWalks"     "League"     "Division"   "PutOuts"    "Assists"    "Errors"
## [19] "Salary"     "NewLeague"
```

```
dim(Hitters)
```

```
## [1] 322 20
```

```
sum(is.na(Hitters$Salary))
```

```
## [1] 59
```

```
#
```

```
Hitters <- na.omit(Hitters)
```

```
dim(Hitters)
```

```
## [1] 263 20
```

```
sum(is.na(Hitters))
```

```
## [1] 0
```

```
library(leaps)
```

```
## Warning: 'leaps' R 4.2.3
```

```
fit <- regsubsets(Salary ~ ., Hitters)
# best regression subset
summary(fit)
```

```
## Subset selection object
## Call: regsubsets.formula(Salary ~ ., Hitters)
## 19 Variables (and intercept)
##           Forced in Forced out
## AtBat      FALSE      FALSE
## Hits       FALSE      FALSE
## HmRun       FALSE      FALSE
## Runs       FALSE      FALSE
## RBI        FALSE      FALSE
## Walks      FALSE      FALSE
## Years      FALSE      FALSE
## CAAtBat    FALSE      FALSE
## CHits      FALSE      FALSE
```

```

## CHmRun      FALSE      FALSE
## CRuns       FALSE      FALSE
## CRBI        FALSE      FALSE
## CWalks      FALSE      FALSE
## LeagueN     FALSE      FALSE
## DivisionW   FALSE      FALSE
## PutOuts     FALSE      FALSE
## Assists     FALSE      FALSE
## Errors      FALSE      FALSE
## NewLeagueN  FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##           AtBat Hits HmRun Runs RBI Walks Years CatBat CHits CHmRun CRuns CRBI
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " "*" " " " " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " "*" " " " " " " " " " " " " " " " " " " "
## 4 ( 1 ) " " "*" " " " " " " " " " " " " " " " " " " "
## 5 ( 1 ) "*" "*" " " " " " " " " " " " " " " " " " " "
## 6 ( 1 ) "*" "*" " " " " " " "*" " " " " " " " " " " " "
## 7 ( 1 ) " " "*" " " " " " " "*" " " "*" "*" "*" " " " "
## 8 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "*" "*" " "
##           CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1 ( 1 ) " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " "*" " " " " "
## 4 ( 1 ) " " " " "*" "*" " " " " " "
## 5 ( 1 ) " " " " "*" "*" " " " " " "
## 6 ( 1 ) " " " " "*" "*" " " " " " "
## 7 ( 1 ) " " " " "*" "*" " " " " " "
## 8 ( 1 ) "*" " " "*" "*" " " " " " "

```

```

sg <- summary(fit)
names(sg)

```

```

## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"

```

```

dim(sg$which)

```

```

## [1] 8 20

```

```

sg$which

```

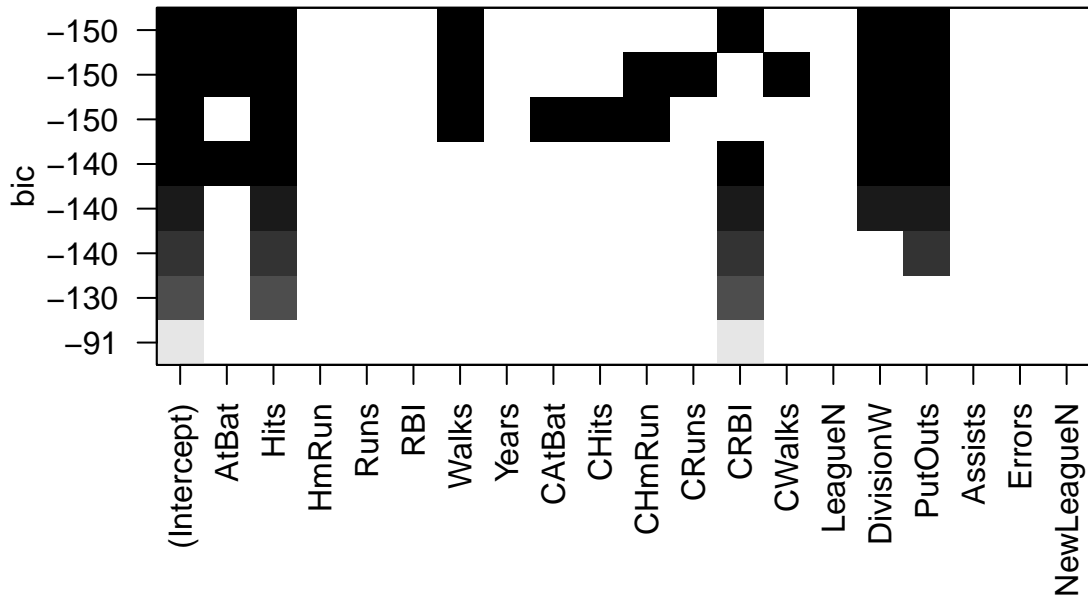
```

## (Intercept) AtBat Hits HmRun Runs RBI Walks Years CatBat CHits CHmRun
## 1 TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 2 TRUE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 3 TRUE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 4 TRUE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 5 TRUE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 6 TRUE TRUE TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE
## 7 TRUE FALSE TRUE FALSE FALSE FALSE TRUE FALSE TRUE TRUE TRUE
## 8 TRUE TRUE TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE TRUE
## CRuns CRBI CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN

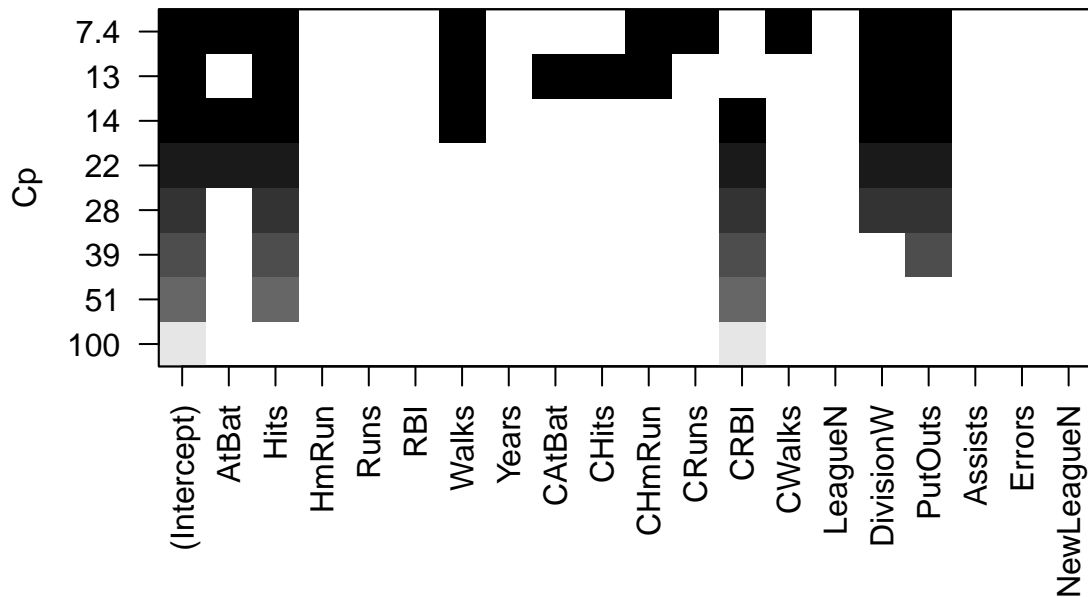
```

```
## 1 FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 2 FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 3 FALSE TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
## 4 FALSE TRUE FALSE FALSE TRUE TRUE FALSE FALSE FALSE
## 5 FALSE TRUE FALSE FALSE TRUE TRUE FALSE FALSE FALSE
## 6 FALSE TRUE FALSE FALSE TRUE TRUE FALSE FALSE FALSE
## 7 FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE
## 8 TRUE FALSE TRUE FALSE TRUE TRUE FALSE FALSE FALSE
```

```
plot(fit) # default: bic
```



```
plot(fit, scale="Cp")
```



```
big <- regsubsets(Salary ~ ., data=Hitters, nvmax=19, nbest=10)
#      8      , 19
#      10
sg <- summary(big)
dim(sg$which)
```

```
## [1] 181 20
```

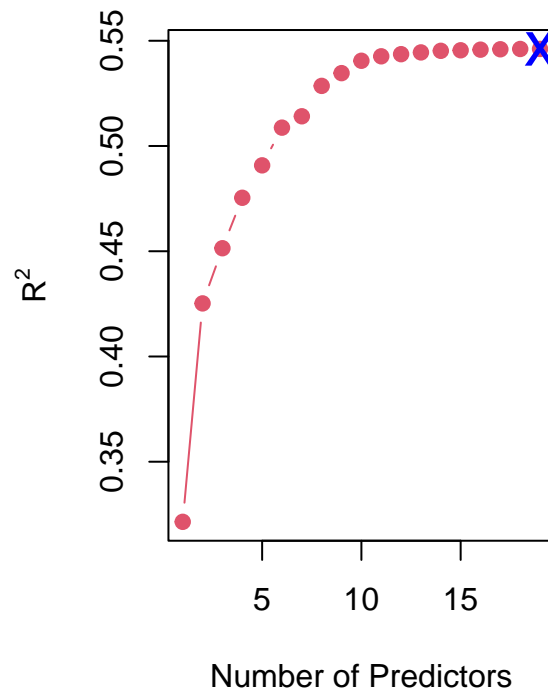
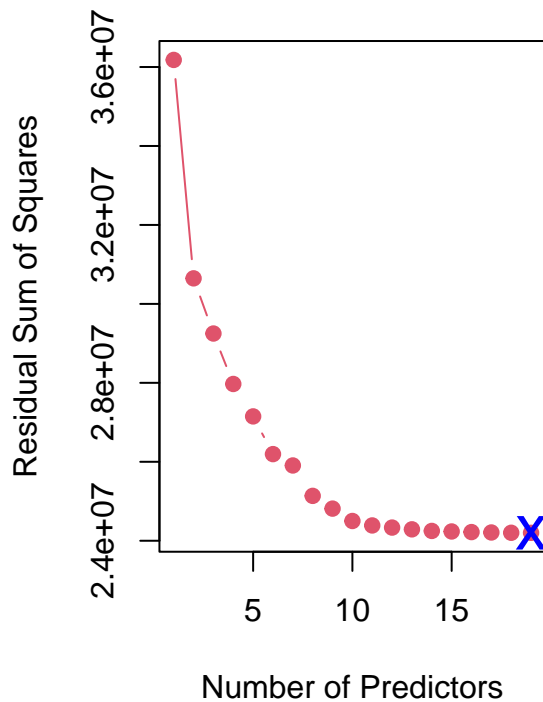
- best : rss, rsq

```
sg.size <- as.numeric(rownames(sg$which))
table(sg.size)
```

```
## sg.size
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
## 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 1
```

```
sg.rss <- tapply(sg$rss, sg.size, min)
w1 <- which.min(sg.rss)
sg.rsq <- tapply(sg$rsq, sg.size, max)
w2 <- which.max(sg.rsq)
par(mfrow=c(1,2))
plot(1:19, sg.rss, type="b", xlab="Number of Predictors",
ylab="Residual Sum of Squares", col=2, pch=19)
```

```
points(w1, sg.rss[w1], pch="x", col="blue", cex=2)
plot(1:19, sg.rsq, type="b", xlab="Number of Predictors",
ylab=expression(R2), col=2, pch=19)
points(w2, sg.rsq[w2], pch="x", col="blue", cex=2)
```



- RSS, R^2
-
-

Forward, Backward selection

```
g.full <- regsubsets(Salary ~., data=Hitters)
g.forw <- regsubsets(Salary ~., data=Hitters, method="forward")
g.back <- regsubsets(Salary ~., data=Hitters, method="backward")
full <- summary(g.full)$which[,-1]
full[full==TRUE] <- 1
forw <- summary(g.forw)$which[,-1]
forw[forw==TRUE] <- 1
back <- summary(g.back)$which[,-1]
back[back==TRUE] <- 1
```

full

##	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAAtBat	CHits	CHmRun	CRuns	CRBI	CWalks
## 1	0	0	0	0	0	0	0	0	0	0	0	1	0
## 2	0	1	0	0	0	0	0	0	0	0	0	1	0
## 3	0	1	0	0	0	0	0	0	0	0	0	1	0
## 4	0	1	0	0	0	0	0	0	0	0	0	1	0
## 5	1	1	0	0	0	0	0	0	0	0	0	1	0
## 6	1	1	0	0	0	1	0	0	0	0	0	1	0
## 7	0	1	0	0	0	1	0	1	1	1	0	0	0
## 8	1	1	0	0	0	1	0	0	0	1	1	0	1
##	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN							
## 1	0		0	0	0	0	0		0				
## 2	0		0	0	0	0	0		0				
## 3	0		0	1	0	0	0		0				
## 4	0		1	1	0	0	0		0				
## 5	0		1	1	0	0	0		0				
## 6	0		1	1	0	0	0		0				
## 7	0		1	1	0	0	0		0				
## 8	0		1	1	0	0	0		0				

forw

## 1	0	0	0	0	0	0	0	0	0	0	0	1	0
## 2	0	1	0	0	0	0	0	0	0	0	0	1	0
## 3	0	1	0	0	0	0	0	0	0	0	0	1	0
## 4	0	1	0	0	0	0	0	0	0	0	0	1	0
## 5	1	1	0	0	0	0	0	0	0	0	0	1	0
## 6	1	1	0	0	0	1	0	0	0	0	0	1	0
## 7	1	1	0	0	0	1	0	0	0	0	0	1	1
## 8	1	1	0	0	0	1	0	0	0	0	1	1	1
##	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN							
## 1	0		0	0	0	0	0		0				
## 2	0		0	0	0	0	0		0				
## 3	0		0	1	0	0	0		0				
## 4	0		1	1	0	0	0		0				
## 5	0		1	1	0	0	0		0				
## 6	0		1	1	0	0	0		0				
## 7	0		1	1	0	0	0		0				
## 8	0		1	1	0	0	0		0				

back

## 1	0	0	0	0	0	0	0	0	0	0	1	0	0
## 2	0	1	0	0	0	0	0	0	0	0	1	0	0
## 3	0	1	0	0	0	0	0	0	0	0	1	0	0
## 4	1	1	0	0	0	0	0	0	0	0	1	0	0
## 5	1	1	0	0	0	1	0	0	0	0	1	0	0
## 6	1	1	0	0	0	1	0	0	0	0	1	0	0
## 7	1	1	0	0	0	1	0	0	0	0	1	0	1


```
## 8      1      1      0      0      0      1      0      0      0      0      1      1      1
## LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1      0      0      0      0      0      0      0
## 2      0      0      0      0      0      0      0
## 3      0      0      1      0      0      0      0
## 4      0      0      1      0      0      0      0
## 5      0      0      1      0      0      0      0
## 6      0      1      1      0      0      0      0
## 7      0      1      1      0      0      0      0
## 8      0      1      1      0      0      0      0
```

```
coef(g.full, 1:5)
```

```
## [[1]]
## (Intercept)          CRBI
## 274.5803864    0.7909536
##
## [[2]]
## (Intercept)      Hits          CRBI
## -47.9559022    3.3008446    0.6898994
##
## [[3]]
## (Intercept)      Hits          CRBI      PutOuts
## -71.4592204    2.8038162    0.6825275    0.2735814
##
## [[4]]
## (Intercept)      Hits          CRBI      DivisionW      PutOuts
## 13.9231044    2.6757978    0.6817790 -139.9538855    0.2735002
##
## [[5]]
## (Intercept)      AtBat          Hits          CRBI      DivisionW      PutOuts
## 97.7684116    -1.4401428    7.1753197    0.6882079 -129.7319386    0.2905164
```

```
coef(g.forw, 1:5)
```

```
## [[1]]
## (Intercept)          CRBI
## 274.5803864    0.7909536
##
## [[2]]
## (Intercept)      Hits          CRBI
## -47.9559022    3.3008446    0.6898994
##
## [[3]]
## (Intercept)      Hits          CRBI      PutOuts
## -71.4592204    2.8038162    0.6825275    0.2735814
##
## [[4]]
## (Intercept)      Hits          CRBI      DivisionW      PutOuts
## 13.9231044    2.6757978    0.6817790 -139.9538855    0.2735002
##
## [[5]]
## (Intercept)      AtBat          Hits          CRBI      DivisionW      PutOuts
```

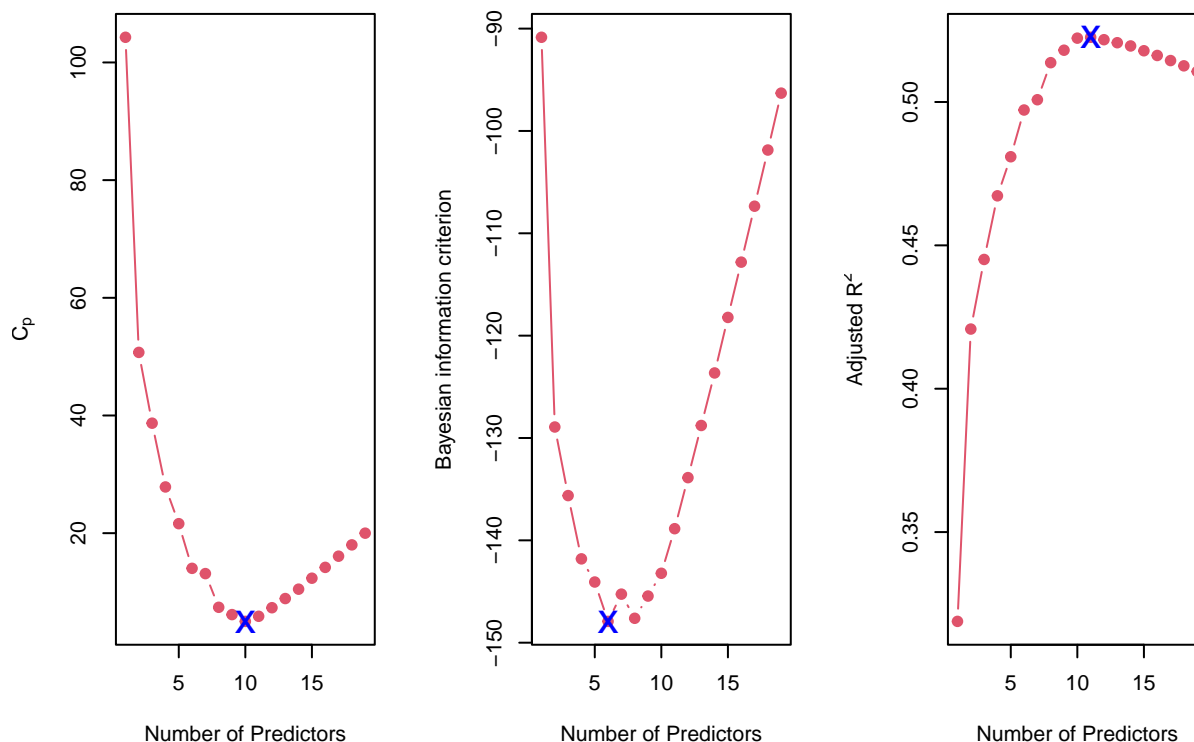
```
## 97.7684116 -1.4401428 7.1753197 0.6882079 -129.7319386 0.2905164
```

```
coef(g.back, 1:5)
```

```
## [[1]]
## (Intercept)      CRuns
## 259.0822757    0.7664116
##
## [[2]]
## (Intercept)      Hits      CRuns
## -50.8174029    3.2257212    0.6614168
##
## [[3]]
## (Intercept)      Hits      CRuns      PutOuts
## -79.3969049    2.6431989    0.6648928    0.3100558
##
## [[4]]
## (Intercept)      AtBat      Hits      CRuns      PutOuts
## 17.5690044   -1.5476983    7.4695315    0.6702041    0.3286804
##
## [[5]]
## (Intercept)      AtBat      Hits      Walks      CRuns      PutOuts
## 16.3605866   -1.9686325    7.8904780    3.6851865    0.6218929    0.3000029
```

Cp(AIC), BIC, adj_R²

```
sg.cp <- tapply(sg$cp, sg.size, min)
w3 <- which.min(sg.cp)
sg.bic <- tapply(sg$bic, sg.size, min)
w4 <- which.min(sg.bic)
sg.adj2 <- tapply(sg$adj2, sg.size, max)
w5 <- which.max(sg.adj2)
par(mfrow=c(1,3))
plot(1:19, sg.cp, type="b", xlab="Number of Predictors",
     ylab=expression(C[p]), col=2, pch=19)
points(w3, sg.cp[w3], pch="x", col="blue", cex=2)
plot(1:19, sg.bic, type="b", xlab="Number of Predictors",
     ylab="Bayesian information criterion", col=2, pch=19)
points(w4, sg.bic[w4], pch="x", col="blue", cex=2)
plot(1:19, sg.adj2, type="b", xlab="Number of Predictors",
     ylab=expression(paste("Adjusted ", R^2)), col=2, pch=19)
points(w5, sg.adj2[w5], pch="x", col="blue", cex=2)
```



```

model1 <- coef(big, which.min(sg$rss))
model2 <- coef(big, which.max(sg$rsq))
model3 <- coef(big, which.max(sg$adjr2))
model4 <- coef(big, which.min(sg$cp))
model5 <- coef(big, which.min(sg$bic))
RES <- matrix(0, 20, 5)
rownames(RES) <- names(model1)
colnames(RES) <- c("rss", "rsq", "adjr2", "cp", "bic")
for (i in 1:5) {
  model <- get(paste("model", i, sep=""))
  w <- match(names(model), rownames(RES))
  RES[w, i] <- model
}
RES

```

##		rss	rsq	adjr2	cp	bic
##	(Intercept)	163.1035878	163.1035878	135.7512195	162.5354420	91.5117981
##	AtBat	-1.9798729	-1.9798729	-2.1277482	-2.1686501	-1.8685892
##	Hits	7.5007675	7.5007675	6.9236994	6.9180175	7.6043976
##	HmRun	4.3308829	4.3308829	0.0000000	0.0000000	0.0000000
##	Runs	-2.3762100	-2.3762100	0.0000000	0.0000000	0.0000000
##	RBI	-1.0449620	-1.0449620	0.0000000	0.0000000	0.0000000
##	Walks	6.2312863	6.2312863	5.6202755	5.7732246	3.6976468
##	Years	-3.4890543	-3.4890543	0.0000000	0.0000000	0.0000000
##	CAtBat	-0.1713405	-0.1713405	-0.1389914	-0.1300798	0.0000000
##	CHits	0.1339910	0.1339910	0.0000000	0.0000000	0.0000000

```
## CHmRun      -0.1728611  -0.1728611   0.0000000   0.0000000   0.0000000
## CRuns       1.4543049   1.4543049   1.4553310   1.4082490   0.0000000
## CRBI        0.8077088   0.8077088   0.7852528   0.7743122   0.6430169
## CWalks      -0.8115709  -0.8115709  -0.8228559  -0.8308264   0.0000000
## LeagueN     62.5994230  62.5994230  43.1116152   0.0000000   0.0000000
## DivisionW  -116.8492456 -116.8492456 -111.1460252 -112.3800575 -122.9515338
## PutOuts     0.2818925   0.2818925   0.2894087   0.2973726   0.2643076
## Assists     0.3710692   0.3710692   0.2688277   0.2831680   0.0000000
## Errors      -3.3607605  -3.3607605   0.0000000   0.0000000   0.0000000
## NewLeagueN  -24.7623251  -24.7623251   0.0000000   0.0000000   0.0000000
```

```
apply(RES, 2, function(t) sum(t!=0)-1)
```

```
##   rss   rsq adjr2    cp   bic
##   19    19    11    10     6
```

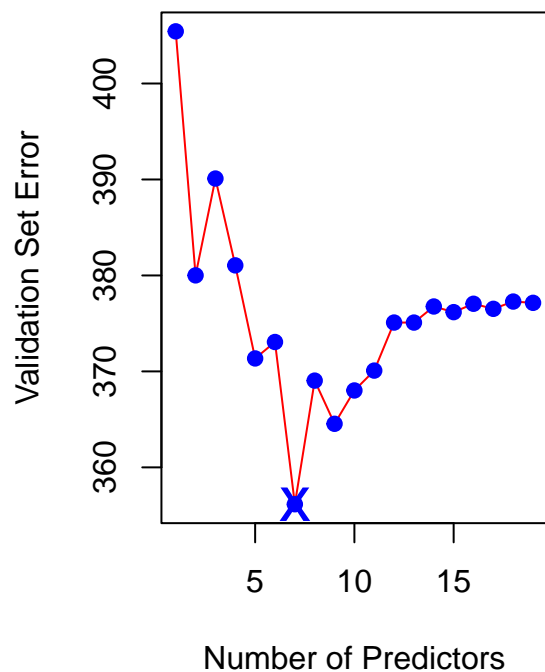
Validation Set

```
set.seed(1)
train <- sample(c(TRUE, FALSE), nrow(Hitters), replace=TRUE)
test <- (!train)
g1 <- regsubsets(Salary ~ ., data=Hitters[train, ], nvmax=19)
test.mat <- model.matrix(Salary~., data=Hitters[test, ])
val.errors <- rep(NA, 19)

for (i in 1:19) {
  coefi <- coef(g1, id=i)
  pred <- test.mat[, names(coefi)] %*% coefi
  val.errors[i] <- sqrt(mean((Hitters$Salary[test]-pred)^2))
}
val.errors
```

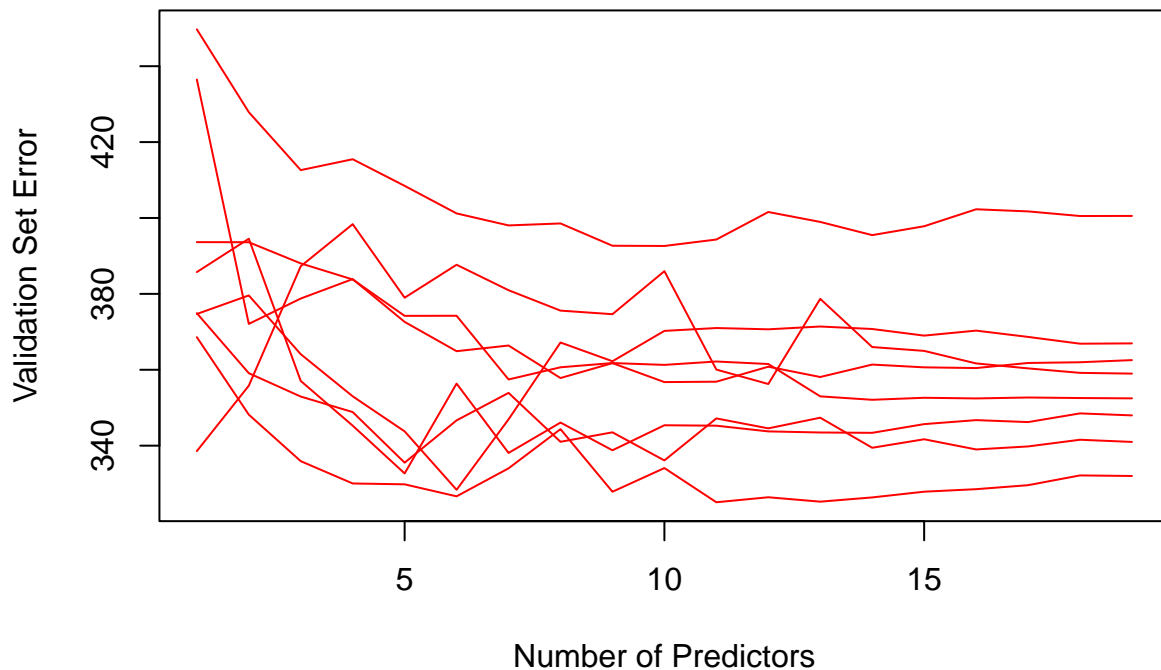
```
## [1] 405.4347 380.0072 390.0970 381.0491 371.3517 373.0627 356.1586 369.0412
## [9] 364.5403 368.0148 370.0855 375.0932 375.0879 376.7641 376.1757 377.0469
## [17] 376.5201 377.2791 377.1447
```

```
w <- which.min(val.errors)
par(mfrow=c(1,2))
plot(1:19, val.errors, type="l", col="red",
     xlab="Number of Predictors", ylab="Validation Set Error")
points(1:19, val.errors, pch=19, col="blue")
points(w, val.errors[w], pch="x", col="blue", cex=2)
```



K-fold validation

```
set.seed(1234)
N <- 8
ERR <- matrix(0, 19, N)
for (k in 1:N) {
  tr <- sample(c(TRUE, FALSE), nrow(Hitters), replace=TRUE)
  tt <- (!tr)
  g <- regsubsets(Salary ~ ., data=Hitters[tr, ], nvmax=19)
  tt.mat <- model.matrix(Salary~., data=Hitters[tt, ])
  for (i in 1:19) {
    coefi <- coef(g, id=i)
    pred <- tt.mat[, names(coefi)] %*% coefi
    ERR[i,k] <- sqrt(mean((Hitters$Salary[tt]-pred)^2))
  }
}
matplot(ERR, type="l", col="red", xlab="Number of Predictors",
lty=1, ylab="Validation Set Error")
```



- Validation Set
- !

```
apply(ERR, 2, which.min)
```

```
## [1] 11  5 14 10 10  6  1  5
```

```
## Define new "predict" function on regsubset
predict.regsubsets <- function(object, newdata, id, ...) {
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id=id)
  xvars <- names(coefi)
  mat[, xvars] %*% coefi
}
set.seed(1)
K <- 10
n <- nrow(Hitters)
fd <- sample(rep(1:K, length=n))
cv.errors <- matrix(NA, n, 19, dimnames=list(NULL, paste(1:19)))

for (i in 1:K) {
  fit <- regsubsets(Salary~., Hitters[fd!=i, ], nvmax=19)
  for (j in 1:19) {
    pred <- predict(fit, Hitters[fd==i, ], id=j)
```

```

    cv.errors[fd==i, j] <- (Hitters$Salary[fd==i]-pred)^2
  }
}

sqrt(apply(cv.errors, 2, mean))

```

```

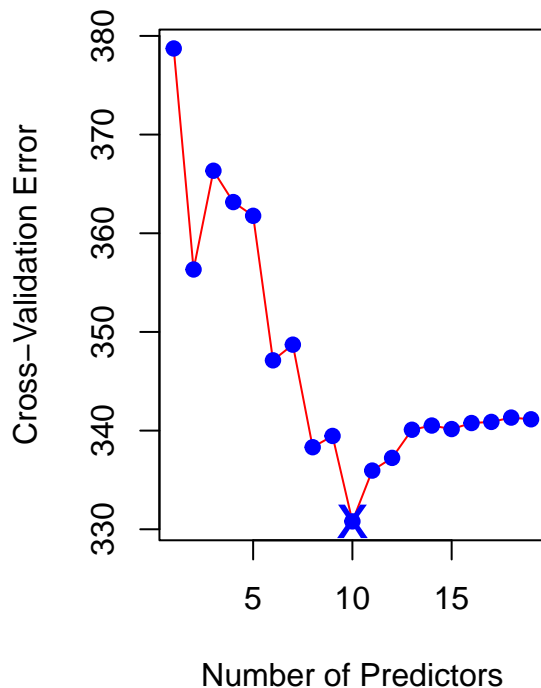
##          1          2          3          4          5          6          7          8
## 378.7273 356.3299 366.3377 363.1608 361.7680 347.1194 348.7096 338.3155
##          9         10         11         12         13         14         15         16
## 339.4583 330.8001 335.9470 337.2333 340.0865 340.5136 340.1557 340.7668
##         17         18         19
## 340.8814 341.3180 341.1447

```

```

K.ERR <- sqrt(apply(cv.errors, 2, mean))
ww <- which.min(K.ERR)
par(mfrow=c(1,2))
plot(1:19, K.ERR, type="l", col="red",
     xlab="Number of Predictors", ylab="Cross-Validation Error")
points(1:19, K.ERR, pch=19, col="blue")
points(ww, K.ERR[ww], pch="x", col="blue", cex=2)

```



```

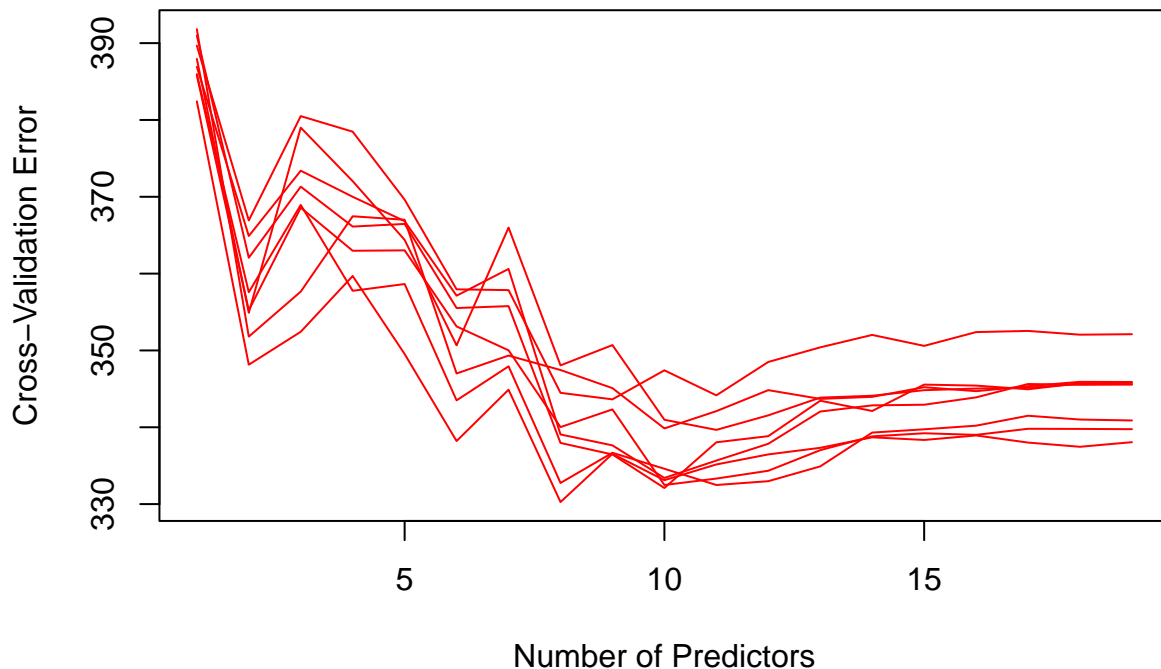
## 10-fold CV with 8 different splits
N <- 8
n <- nrow(Hitters)

```

```

ERR <- matrix(0, 19, N)
set.seed(1234)
for (k in 1:N) {
  fd <- sample(rep(1:K, length=n))
  CVR <- matrix(NA, n, 19)
  for (i in 1:K) {
    f <- regsubsets(Salary~., data=Hitters[fd!=i, ], nvmax=19)
    for (j in 1:19) {
      pred <- predict(f, Hitters[fd==i, ], id=j)
      CVR[fd==i, j] <- (Hitters$Salary[fd==i]-pred)^2
    }
  }
  ERR[,k] <- sqrt(apply(CVR, 2, mean))
}
matplot(ERR, type="l", col="red", xlab="Number of Predictors",
lty=1, ylab="Cross-Validation Error")

```



```

apply(ERR, 2, which.min)

```

```

## [1] 10 10 10 9 11 10 11 8

```

```

set.seed(111)
fd <- sample(rep(1:K, length=n))
CVR.1se <- matrix(NA, n, 19)

```



```

for (i in 1:K) {
  fit <- regsubsets(Salary~., Hitters[fd!=i, ], nvmax=19)
  for (j in 1:19) {
    pred <- predict(fit, Hitters[fd==i, ], id=j)
    CVR.1se[fd==i, j] <- Hitters$Salary[fd==i]-pred
  }
}
avg <- sqrt(apply(CVR.1se^2, 2, mean))
se <- apply(CVR.1se, 2, sd)/sqrt(n)
PE <- cbind(avg - se, avg, avg + se)
data.frame(lwr=PE[,1], mean=PE[,2], upp=PE[,3])

```

```

##          lwr      mean      upp
## 1  369.1514 393.4582 417.7651
## 2  352.7224 375.9457 399.1690
## 3  356.6445 380.1251 403.6058
## 4  348.5281 371.4749 394.4218
## 5  350.8331 373.9322 397.0313
## 6  339.3305 361.6738 384.0170
## 7  345.5366 368.2846 391.0326
## 8  330.4133 352.1664 373.9196
## 9  329.9209 351.6399 373.3590
## 10 317.9704 338.9042 359.8380
## 11 322.5224 343.7582 364.9941
## 12 324.6231 345.9976 367.3720
## 13 327.7008 349.2755 370.8502
## 14 327.2778 348.8254 370.3730
## 15 328.2571 349.8688 371.4806
## 16 328.8777 350.5306 372.1836
## 17 327.9254 349.5151 371.1048
## 18 328.4501 350.0743 371.6985
## 19 328.5078 350.1358 371.7638

```

```
which.min(PE[,2])
```

```
## [1] 10
```

```

w <- which.min(PE[,2])
which(PE[w, 1] < PE[,2] & PE[w, 3] > PE[,2])

```

```
## [1] 8 9 10 11 12 13 14 15 16 17 18 19
```

```
min(which(PE[w, 1] < PE[,2] & PE[w, 3] > PE[,2]))
```

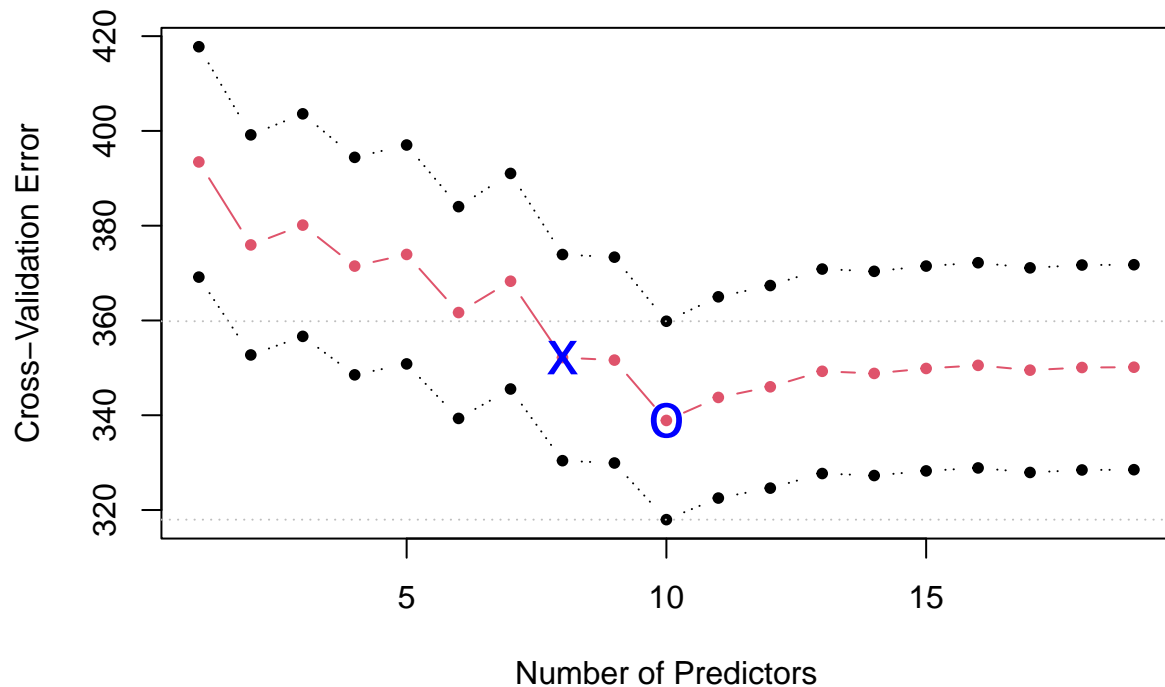
```
## [1] 8
```

```

matplot(1:19, PE, type="b", col=c(1,2,1), lty=c(3,1,3), pch=20,
xlab="Number of Predictors", ylab="Cross-Validation Error")
abline(h=PE[w, 1], lty=3, col="gray")
abline(h=PE[w, 3], lty=3, col="gray")

```

```
points(which.min(avg), PE[which.min(avg),2],
pch="o",col="blue",cex=2)
up <- which(PE[,2] < PE[which.min(PE[,2]),3])
points(min(up), PE[min(up),2], pch="x", col="blue", cex=2)
```



- CV Error 10 upper, lower 8 .

Shrinkage Methods

- shrinkage : , ,
- Ridge: $L2$ -norm penalty
 $-\lambda$ $lambda$.
- Lasso (Least absolute shrinkage and selection operator): $L1$ -norm penalty

```
library(glmnet)
```

1. Ridge regression

```
##      : Matrix
```

```
## Loaded glmnet 4.1-8
```

```
x0 <- model.matrix(Salary~., Hitters)[, -1]
y <- Hitters$Salary

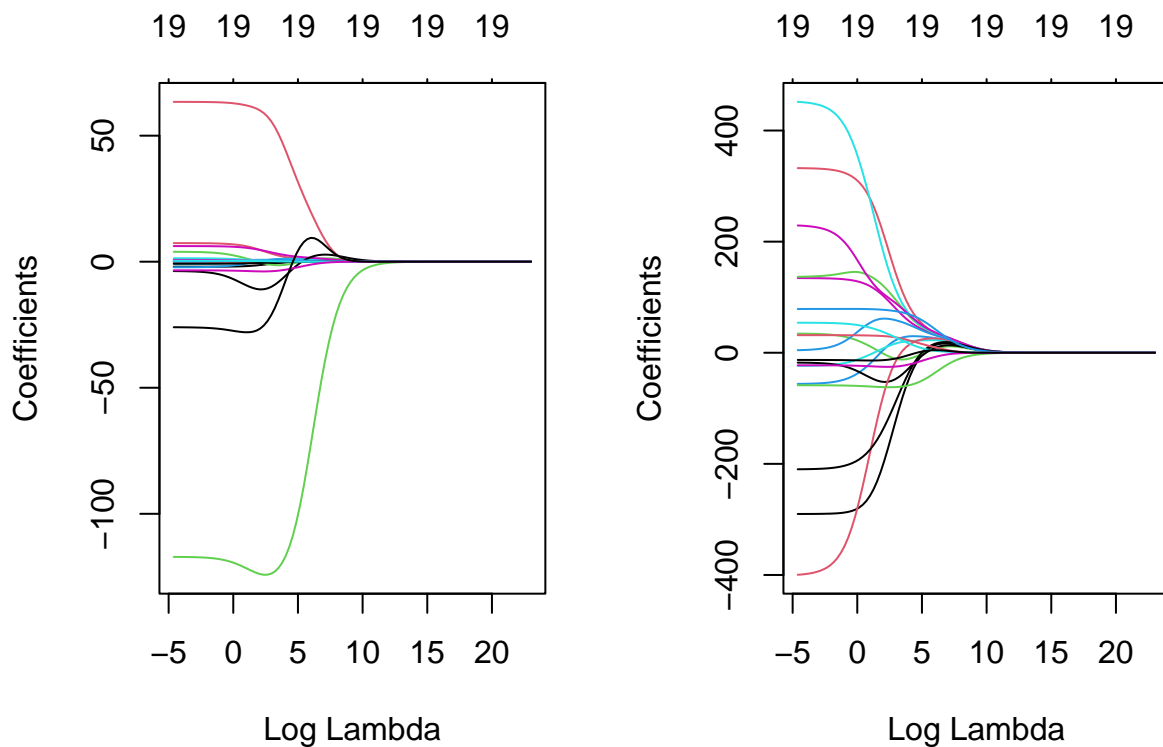
grid <- 10^seq(10, -2, length=100)
g1 <- glmnet(x0, y, alpha=0, lambda=grid)

par(mfrow=c(1,2))
plot(g1, "lambda", label=TRUE)

fun <- function(t) sqrt(var(t)*(length(t)-1)/length(t))
sdx <- matrix(apply(x0, 2, fun), dim(x0)[2], dim(x0)[1])

x <- x0/t(sdx)
g2 <- glmnet(x, y, alpha=0, lambda=grid)

plot(g2, "lambda", label=TRUE)
```



- glmnet
- 19 shrinkage method (= 0)

```
data.frame(sd_g1=apply(x0, 2, sd), sd_g2=apply(x, 2, sd))
```

```
##          sd_g1      sd_g2
```

```
## AtBat      147.3072088 1.001907
## Hits       45.1253259 1.001907
## HmRun       8.7571077 1.001907
## Runs       25.5398156 1.001907
## RBI        25.8827143 1.001907
## Walks      21.7180561 1.001907
## Years       4.7936159 1.001907
## CAtBat     2286.5829295 1.001907
## CHits      648.1996437 1.001907
## CHmRun     82.1975815 1.001907
## CRuns      331.1985706 1.001907
## CRBI       323.3676682 1.001907
## CWalks     264.0558680 1.001907
## LeagueN     0.5001378 1.001907
## DivisionW   0.5008628 1.001907
## PutOuts    279.9345755 1.001907
## Assists    145.0805766 1.001907
## Errors      6.6065742 1.001907
## NewLeagueN  0.4996443 1.001907
```

```
names(g2)
```

```
## [1] "a0"      "beta"    "df"      "dim"     "lambda"  "dev.ratio"
## [7] "nulldev" "npasses" "jerr"    "offset"  "call"    "nobs"
```

```
data.frame(lambda=g2$lambda, df=g2$df)
```

```
##      lambda df
## 1  1.000000e+10 19
## 2  7.564633e+09 19
## 3  5.722368e+09 19
## 4  4.328761e+09 19
## 5  3.274549e+09 19
## 6  2.477076e+09 19
## 7  1.873817e+09 19
## 8  1.417474e+09 19
## 9  1.072267e+09 19
## 10 8.111308e+08 19
## 11 6.135907e+08 19
## 12 4.641589e+08 19
## 13 3.511192e+08 19
## 14 2.656088e+08 19
## 15 2.009233e+08 19
## 16 1.519911e+08 19
## 17 1.149757e+08 19
## 18 8.697490e+07 19
## 19 6.579332e+07 19
## 20 4.977024e+07 19
## 21 3.764936e+07 19
## 22 2.848036e+07 19
## 23 2.154435e+07 19
## 24 1.629751e+07 19
## 25 1.232847e+07 19
```

26 9.326033e+06 19
27 7.054802e+06 19
28 5.336699e+06 19
29 4.037017e+06 19
30 3.053856e+06 19
31 2.310130e+06 19
32 1.747528e+06 19
33 1.321941e+06 19
34 1.000000e+06 19
35 7.564633e+05 19
36 5.722368e+05 19
37 4.328761e+05 19
38 3.274549e+05 19
39 2.477076e+05 19
40 1.873817e+05 19
41 1.417474e+05 19
42 1.072267e+05 19
43 8.111308e+04 19
44 6.135907e+04 19
45 4.641589e+04 19
46 3.511192e+04 19
47 2.656088e+04 19
48 2.009233e+04 19
49 1.519911e+04 19
50 1.149757e+04 19
51 8.697490e+03 19
52 6.579332e+03 19
53 4.977024e+03 19
54 3.764936e+03 19
55 2.848036e+03 19
56 2.154435e+03 19
57 1.629751e+03 19
58 1.232847e+03 19
59 9.326033e+02 19
60 7.054802e+02 19
61 5.336699e+02 19
62 4.037017e+02 19
63 3.053856e+02 19
64 2.310130e+02 19
65 1.747528e+02 19
66 1.321941e+02 19
67 1.000000e+02 19
68 7.564633e+01 19
69 5.722368e+01 19
70 4.328761e+01 19
71 3.274549e+01 19
72 2.477076e+01 19
73 1.873817e+01 19
74 1.417474e+01 19
75 1.072267e+01 19
76 8.111308e+00 19
77 6.135907e+00 19
78 4.641589e+00 19
79 3.511192e+00 19

```
## 80 2.656088e+00 19
## 81 2.009233e+00 19
## 82 1.519911e+00 19
## 83 1.149757e+00 19
## 84 8.697490e-01 19
## 85 6.579332e-01 19
## 86 4.977024e-01 19
## 87 3.764936e-01 19
## 88 2.848036e-01 19
## 89 2.154435e-01 19
## 90 1.629751e-01 19
## 91 1.232847e-01 19
## 92 9.326033e-02 19
## 93 7.054802e-02 19
## 94 5.336699e-02 19
## 95 4.037017e-02 19
## 96 3.053856e-02 19
## 97 2.310130e-02 19
## 98 1.747528e-02 19
## 99 1.321941e-02 19
## 100 1.000000e-02 19
```

```
data.frame(log.lambda=round(log(g2$lambda), 4), df=g2$df)
```

```
##      log.lambda df
## 1      23.0259 19
## 2      22.7467 19
## 3      22.4676 19
## 4      22.1885 19
## 5      21.9094 19
## 6      21.6303 19
## 7      21.3512 19
## 8      21.0721 19
## 9      20.7930 19
## 10     20.5139 19
## 11     20.2348 19
## 12     19.9557 19
## 13     19.6766 19
## 14     19.3975 19
## 15     19.1184 19
## 16     18.8393 19
## 17     18.5602 19
## 18     18.2811 19
## 19     18.0020 19
## 20     17.7229 19
## 21     17.4438 19
## 22     17.1647 19
## 23     16.8856 19
## 24     16.6065 19
## 25     16.3274 19
## 26     16.0483 19
## 27     15.7692 19
## 28     15.4901 19
## 29     15.2110 19
```

## 30	14.9319	19
## 31	14.6528	19
## 32	14.3737	19
## 33	14.0946	19
## 34	13.8155	19
## 35	13.5364	19
## 36	13.2573	19
## 37	12.9782	19
## 38	12.6991	19
## 39	12.4200	19
## 40	12.1409	19
## 41	11.8618	19
## 42	11.5827	19
## 43	11.3036	19
## 44	11.0245	19
## 45	10.7454	19
## 46	10.4663	19
## 47	10.1872	19
## 48	9.9081	19
## 49	9.6290	19
## 50	9.3499	19
## 51	9.0708	19
## 52	8.7917	19
## 53	8.5126	19
## 54	8.2335	19
## 55	7.9544	19
## 56	7.6753	19
## 57	7.3962	19
## 58	7.1171	19
## 59	6.8380	19
## 60	6.5589	19
## 61	6.2798	19
## 62	6.0007	19
## 63	5.7216	19
## 64	5.4425	19
## 65	5.1634	19
## 66	4.8843	19
## 67	4.6052	19
## 68	4.3261	19
## 69	4.0470	19
## 70	3.7679	19
## 71	3.4888	19
## 72	3.2097	19
## 73	2.9306	19
## 74	2.6515	19
## 75	2.3724	19
## 76	2.0933	19
## 77	1.8142	19
## 78	1.5351	19
## 79	1.2560	19
## 80	0.9769	19
## 81	0.6978	19
## 82	0.4187	19
## 83	0.1396	19

```
## 84      -0.1396 19
## 85      -0.4187 19
## 86      -0.6978 19
## 87      -0.9769 19
## 88      -1.2560 19
## 89      -1.5351 19
## 90      -1.8142 19
## 91      -2.0933 19
## 92      -2.3724 19
## 93      -2.6515 19
## 94      -2.9306 19
## 95      -3.2097 19
## 96      -3.4888 19
## 97      -3.7679 19
## 98      -4.0470 19
## 99      -4.3261 19
## 100     -4.6052 19
```

```
dim(coef(g2))
```

```
## [1] 20 100
```

```
coef(g2)[, c("s0", "s10", "s20", "s30")]
```

```
## 20 x 4 sparse Matrix of class "dgCMatrix"
##              s0              s10              s20              s30
## (Intercept)  5.359257e+02  5.359228e+02  5.358758e+02  5.351104e+02
## AtBat        8.003360e-06  1.304343e-04  2.125620e-03  3.460670e-02
## Hits        8.893439e-06  1.449404e-04  2.362028e-03  3.845827e-02
## HmRun        6.954354e-06  1.133381e-04  1.846993e-03  3.006538e-02
## Runs        8.511971e-06  1.387233e-04  2.260696e-03  3.680432e-02
## RBI          9.112034e-06  1.485028e-04  2.420054e-03  3.939536e-02
## Walks        8.998709e-06  1.466559e-04  2.389972e-03  3.890955e-02
## Years        8.122690e-06  1.323790e-04  2.157288e-03  3.511605e-02
## CAtBat       1.066656e-05  1.738376e-04  2.832917e-03  4.611551e-02
## CHits        1.112828e-05  1.813623e-04  2.955542e-03  4.811156e-02
## CHmRun       1.064214e-05  1.734395e-04  2.826419e-03  4.600702e-02
## CRuns        1.140740e-05  1.859114e-04  3.029659e-03  4.931420e-02
## CRBI         1.149433e-05  1.873281e-04  3.052739e-03  4.968778e-02
## CWalks       9.930368e-06  1.618394e-04  2.637348e-03  4.292109e-02
## LeagueN     -2.895411e-07  -4.718666e-06  -7.686756e-05  -1.243483e-03
## DivisionW   -3.902926e-06  -6.360786e-05  -1.036620e-03  -1.688637e-02
## PutOuts      6.091765e-06  9.928031e-05  1.617947e-03  2.634854e-02
## Assists      5.156775e-07  8.404231e-06  1.369608e-04  2.230188e-03
## Errors      -1.094908e-07  -1.784441e-06  -2.908560e-05  -4.749957e-04
## NewLeagueN  -5.746384e-08  -9.364259e-07  -1.523710e-05  -2.418792e-04
```

```
g2$lambda[50]
```

```
## [1] 11497.57
```



```
coef(g2)[,50]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 407.35605020  5.43369951  6.22356733  4.58549860  5.88086202  6.19593490
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##  6.27798659  5.29979792  7.14752789  7.53950891  7.18234860  7.72864834
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##  7.79069879  6.59289871  0.04244485 -3.10715882  4.60526283  0.37837242
##      Errors      NewLeagueN
## -0.13519479  0.15032295
```

```
sqrt(sum(coef(g2)[-1, 50]^2))
```

```
## [1] 24.17063
```

```
g2$lambda[60]
```

```
## [1] 705.4802
```

```
coef(g2)[,60]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
##  54.325200  16.483353  29.555975  10.312055  23.903039  21.885732
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##  28.610668  12.422480  24.725984  30.242802  27.711305  30.926412
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##  31.566475  18.948465  6.830714 -27.324449  33.115519  2.325614
##      Errors      NewLeagueN
## -4.639451  4.294655
```

```
sqrt(sum(coef(g2)[-1, 60]^2))
```

```
## [1] 99.07558
```

```
l2norm <- function(t) sqrt(sum(t^2))
l2 <- apply(coef(g2)[-1,], 2, l2norm)
data.frame(log_lambda=round(log(g2$lambda), 4),
l2norm=round(l2, 4))
```

```
##      log_lambda  l2norm
## s0      23.0259  0.0000
## s1      22.7467  0.0000
## s2      22.4676  0.0001
## s3      22.1885  0.0001
## s4      21.9094  0.0001
## s5      21.6303  0.0001
## s6      21.3512  0.0002
## s7      21.0721  0.0003
## s8      20.7930  0.0003
```

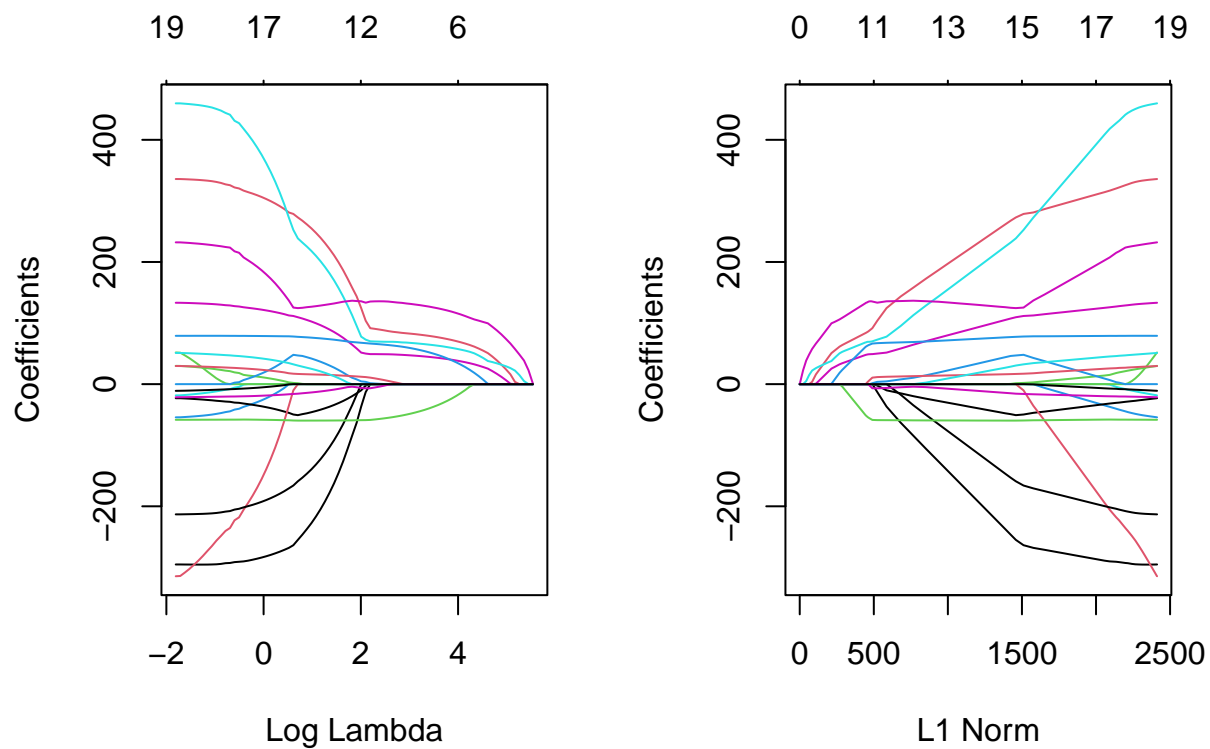
## s9	20.5139	0.0004
## s10	20.2348	0.0006
## s11	19.9557	0.0008
## s12	19.6766	0.0010
## s13	19.3975	0.0013
## s14	19.1184	0.0018
## s15	18.8393	0.0023
## s16	18.5602	0.0031
## s17	18.2811	0.0041
## s18	18.0020	0.0054
## s19	17.7229	0.0071
## s20	17.4438	0.0094
## s21	17.1647	0.0125
## s22	16.8856	0.0165
## s23	16.6065	0.0218
## s24	16.3274	0.0288
## s25	16.0483	0.0380
## s26	15.7692	0.0503
## s27	15.4901	0.0664
## s28	15.2110	0.0878
## s29	14.9319	0.1161
## s30	14.6528	0.1534
## s31	14.3737	0.2027
## s32	14.0946	0.2678
## s33	13.8155	0.3537
## s34	13.5364	0.4672
## s35	13.2573	0.6168
## s36	12.9782	0.8141
## s37	12.6991	1.0739
## s38	12.4200	1.4157
## s39	12.1409	1.8613
## s40	11.8618	2.4473
## s41	11.5827	3.2124
## s42	11.3036	4.2074
## s43	11.0245	5.4950
## s44	10.7454	7.1502
## s45	10.4663	9.2606
## s46	10.1872	11.9229
## s47	9.9081	15.2373
## s48	9.6290	19.2964
## s49	9.3499	24.1706
## s50	9.0708	29.8882
## s51	8.7917	36.4179
## s52	8.5126	43.6571
## s53	8.2335	51.4379
## s54	7.9544	59.5465
## s55	7.6753	67.7614
## s56	7.3962	75.8939
## s57	7.1171	83.8348
## s58	6.8380	91.5394
## s59	6.5589	99.0756
## s60	6.2798	106.5535
## s61	6.0007	114.1538
## s62	5.7216	122.1068

```
## s63      5.4425 130.7626
## s64      5.1634 140.4553
## s65      4.8843 151.7662
## s66      4.6052 165.3257
## s67      4.3261 181.6929
## s68      4.0470 201.4140
## s69      3.7679 224.9535
## s70      3.4888 252.3551
## s71      3.2097 283.4422
## s72      2.9306 317.9353
## s73      2.6515 355.1502
## s74      2.3724 394.3546
## s75      2.0933 434.7652
## s76      1.8142 475.6017
## s77      1.5351 516.2080
## s78      1.2560 555.8559
## s79      0.9769 594.1277
## s80      0.6978 629.9067
## s81      0.4187 663.6317
## s82      0.1396 694.1000
## s83     -0.1396 721.6804
## s84     -0.4187 745.5967
## s85     -0.6978 766.0221
## s86     -0.9769 782.8521
## s87     -1.2560 796.5702
## s88     -1.5351 807.6860
## s89     -1.8142 816.3247
## s90     -2.0933 823.0769
## s91     -2.3724 828.2603
## s92     -2.6515 832.1470
## s93     -2.9306 835.2884
## s94     -3.2097 837.5397
## s95     -3.4888 839.4284
## s96     -3.7679 840.7913
## s97     -4.0470 841.8099
## s98     -4.3261 842.6548
## s99     -4.6052 843.1611
```

2. Lasso

- Ridge regression subset selection .
- Ridge 0 .
- Lasso 0
- 0 .
- optimal model optimal lambda .

```
g3 <- glmnet(x, y, alpha=1)
par(mfrow=c(1,2))
plot(g3, "lambda", label=TRUE)
plot(g3, "norm", label=TRUE)
```



• lambda 0 .

```
dim(coef(g3))
```

```
## [1] 20 80
```

```
coef(g3)[, c("s0", "s10", "s40", "s60")]
```

```
## 20 x 4 sparse Matrix of class "dgCMatrix"
##           s0          s10          s40          s60
## (Intercept) 535.9259 222.0974180 40.652951 153.262451
## AtBat      .          .          -84.358820 -283.848789
## Hits       .          50.8499527 158.005411 306.337268
## HmRun      .          .          .          11.386077
## Runs       .          .          .          -26.445699
## RBI        .          .          .          .
## Walks      .          25.4064700 64.595375 121.943910
## Years      .          .          -19.691966 -36.282468
## CAtBat     .          .          .          -155.199102
## CHits      .          .          .          .
## CHmRun     .          .          10.478327 17.456030
## CRuns      .          37.9218169 112.825764 375.328259
## CRBI       .          99.5844116 136.599052 186.323253
## CWalks     .          .          -34.517519 -192.606975
```

```
## LeagueN      .      .      13.103378    23.517927
## DivisionW    .      .      -59.253067   -58.306307
## PutOuts      .      0.4926175  68.973044    78.789863
## Assists      .      .           .         42.115015
## Errors       .      .      -4.200133   -18.955204
## NewLeagueN   .      .           .        -5.314774
```

```
data.frame(lambda=g3$lambda, df=g3$df)
```

```
##      lambda df
## 1 255.2820965 0
## 2 232.6035386 1
## 3 211.9396813 2
## 4 193.1115442 2
## 5 175.9560468 3
## 6 160.3245966 4
## 7 146.0818013 4
## 8 133.1042967 4
## 9 121.2796778 4
## 10 110.5055255 4
## 11 100.6885192 5
## 12 91.7436287 5
## 13 83.5933775 5
## 14 76.1671723 5
## 15 69.4006906 6
## 16 63.2353245 6
## 17 57.6176726 6
## 18 52.4990774 6
## 19 47.8352040 6
## 20 43.5856563 6
## 21 39.7136268 6
## 22 36.1855776 6
## 23 32.9709506 6
## 24 30.0419022 6
## 25 27.3730624 6
## 26 24.9413150 6
## 27 22.7255973 6
## 28 20.7067179 6
## 29 18.8671902 6
## 30 17.1910810 7
## 31 15.6638727 7
## 32 14.2723374 7
## 33 13.0044223 9
## 34 11.8491453 9
## 35 10.7964999 9
## 36 9.8373686 9
## 37 8.9634439 9
## 38 8.1671562 11
## 39 7.4416086 11
## 40 6.7805166 12
## 41 6.1781542 12
## 42 5.6293040 13
## 43 5.1292121 13
## 44 4.6735471 13
```

```
## 45 4.2583620 13
## 46 3.8800609 13
## 47 3.5353670 13
## 48 3.2212947 13
## 49 2.9351238 13
## 50 2.6743755 13
## 51 2.4367913 13
## 52 2.2203135 14
## 53 2.0230670 15
## 54 1.8433433 15
## 55 1.6795857 17
## 56 1.5303760 17
## 57 1.3944216 17
## 58 1.2705450 17
## 59 1.1576733 17
## 60 1.0548288 17
## 61 0.9611207 17
## 62 0.8757374 17
## 63 0.7979393 17
## 64 0.7270526 17
## 65 0.6624632 18
## 66 0.6036118 18
## 67 0.5499886 18
## 68 0.5011291 17
## 69 0.4566102 18
## 70 0.4160462 18
## 71 0.3790858 18
## 72 0.3454089 18
## 73 0.3147237 18
## 74 0.2867645 18
## 75 0.2612891 18
## 76 0.2380769 18
## 77 0.2169268 18
## 78 0.1976557 18
## 79 0.1800965 18
## 80 0.1640972 19
```

```
dim(g3$beta)
```

```
## [1] 19 80
```

```
df2 <- apply(g3$beta, 2, function(t) sum(t!=0))
```

```
data.frame(df1=g3$df, df2=df2)
```

```
##      df1 df2
## s0     0  0
## s1     1  1
## s2     2  2
## s3     2  2
## s4     3  3
## s5     4  4
```

## s6	4	4
## s7	4	4
## s8	4	4
## s9	4	4
## s10	5	5
## s11	5	5
## s12	5	5
## s13	5	5
## s14	6	6
## s15	6	6
## s16	6	6
## s17	6	6
## s18	6	6
## s19	6	6
## s20	6	6
## s21	6	6
## s22	6	6
## s23	6	6
## s24	6	6
## s25	6	6
## s26	6	6
## s27	6	6
## s28	6	6
## s29	7	7
## s30	7	7
## s31	7	7
## s32	9	9
## s33	9	9
## s34	9	9
## s35	9	9
## s36	9	9
## s37	11	11
## s38	11	11
## s39	12	12
## s40	12	12
## s41	13	13
## s42	13	13
## s43	13	13
## s44	13	13
## s45	13	13
## s46	13	13
## s47	13	13
## s48	13	13
## s49	13	13
## s50	13	13
## s51	14	14
## s52	15	15
## s53	15	15
## s54	17	17
## s55	17	17
## s56	17	17
## s57	17	17
## s58	17	17
## s59	17	17

```
## s60 17 17
## s61 17 17
## s62 17 17
## s63 17 17
## s64 18 18
## s65 18 18
## s66 18 18
## s67 17 17
## s68 18 18
## s69 18 18
## s70 18 18
## s71 18 18
## s72 18 18
## s73 18 18
## s74 18 18
## s75 18 18
## s76 18 18
## s77 18 18
## s78 18 18
## s79 19 19
```

Selecting the Tuning Parameter in shrinkage methods

```
set.seed(123)
train <- sample(1:nrow(x), nrow(x)/2)
test <- (-train)
y.test <- y[test]
grid <- 10^seq(10, -2, length=100)
r1 <- glmnet(x[train, ], y[train], alpha=0, lambda=grid)
ss <- 0:(length(r1$lambda)-1)
Err <- NULL
for (i in 1:length(r1$lambda)) {
  r1.pred <- predict(r1, s=ss[i], newx=x[test, ])
  Err[i] <- mean((r1.pred - y.test)^2)
}
wh <- which.min(Err)
lam.opt <- r1$lambda[wh]
```

```
r.full <- glmnet(x, y, alpha=0, lambda=grid)
r.full$beta[,wh]
```

One standard error rule

##	AtBat	Hits	HmRun	Runs	RBI	Walks	Years
##	-290.21072	332.28861	34.40765	-56.04846	-23.66939	134.40389	-17.76977
##	CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	-399.62021	136.72188	4.61810	451.57447	229.04284	-209.74151	31.65086
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN		
##	-58.53064	78.79843	54.03884	-22.57788	-12.96145		


```
predict(r.full, type="coefficients", s=lam.opt)
```

```
## 20 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 164.11322
## AtBat      -290.21072
## Hits       332.28861
## HmRun       34.40765
## Runs       -56.04846
## RBI        -23.66939
## Walks      134.40389
## Years      -17.76977
## CAtBat     -399.62021
## CHits      136.72188
## CHmRun       4.61810
## CRuns      451.57447
## CRBI       229.04284
## CWalks     -209.74151
## LeagueN     31.65086
## DivisionW   -58.53064
## PutOuts     78.79843
## Assists     54.03884
## Errors     -22.57788
## NewLeagueN -12.96145
```

- validation set CV .

```
set.seed(1)
cv.r <- cv.glmnet(x, y, alpha=0, nfolds=10)
names(cv.r)
```

Cross-Validation(One standard error rule)

```
## [1] "lambda"      "cvm"          "cvstd"        "cvup"         "cvlo"
## [6] "nzero"       "call"         "name"         "glmnet.fit"   "lambda.min"
## [11] "lambda.1se"  "index"
```

```
cbind(cv.r$cvlo, cv.r$cvm, cv.r$cvup)
```

```
##           [,1]      [,2]      [,3]
## [1,] 190140.46 203414.3 216688.2
## [2,] 188435.90 201753.4 215070.9
## [3,] 187999.83 201122.4 214244.9
## [4,] 187767.54 200881.5 213995.6
## [5,] 187513.77 200618.5 213723.2
## [6,] 187236.65 200331.2 213425.7
## [7,] 186934.18 200017.6 213101.1
## [8,] 186604.21 199675.6 212747.0
## [9,] 186244.44 199302.8 212361.1
```

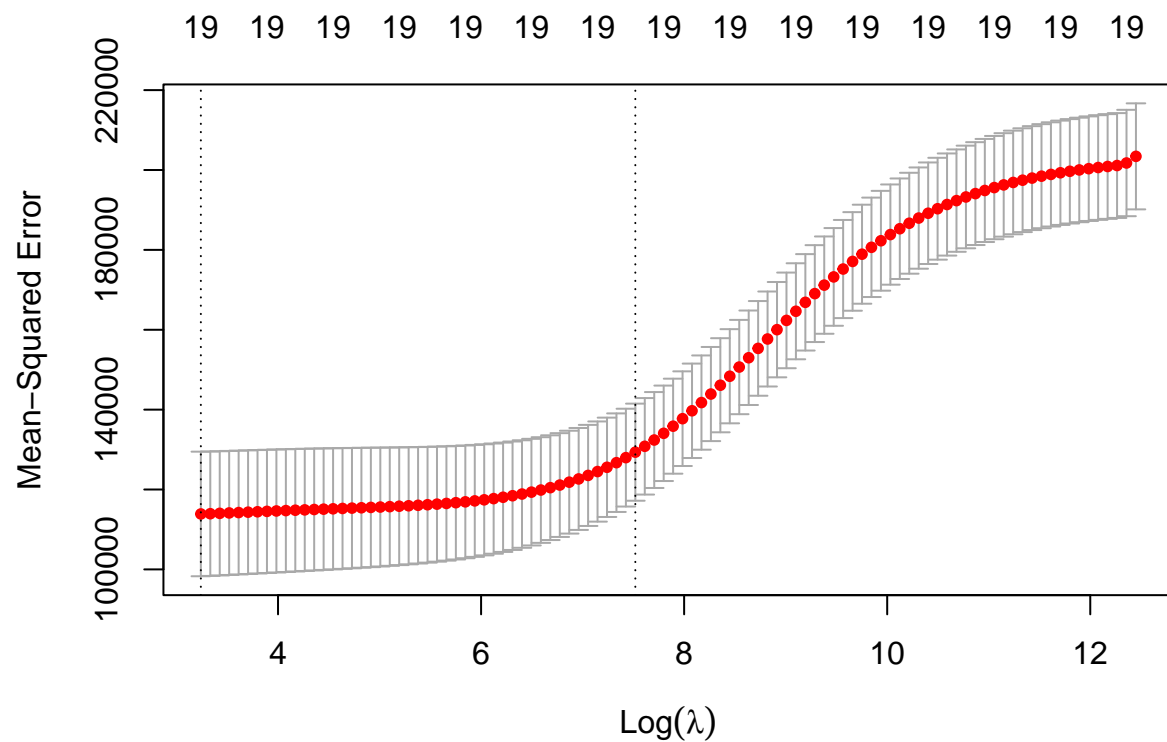
```

## [10,] 185852.45 198896.5 211940.6
## [11,] 185425.62 198454.2 211482.8
## [12,] 184961.20 197973.0 210984.8
## [13,] 184456.31 197449.9 210443.5
## [14,] 183907.89 196881.8 209855.6
## [15,] 183312.74 196265.3 209217.9
## [16,] 182667.47 195597.1 208526.6
## [17,] 181968.87 194873.7 207778.5
## [18,] 181213.39 194091.5 206969.6
## [19,] 180397.47 193247.0 206096.5
## [20,] 179517.56 192336.4 205155.2
## [21,] 178570.15 191356.2 204142.2
## [22,] 177551.80 190302.9 203053.9
## [23,] 176459.22 189173.1 201887.0
## [24,] 175289.34 187963.8 200638.3
## [25,] 174039.38 186672.2 199305.0
## [26,] 172706.96 185296.0 197885.0
## [27,] 171290.15 183833.2 196376.3
## [28,] 169787.63 182282.8 194778.0
## [29,] 168198.74 180644.2 193089.7
## [30,] 166523.63 178917.8 191311.9
## [31,] 164763.34 177104.9 189446.4
## [32,] 162919.90 175207.9 187495.8
## [33,] 160996.41 173230.3 185464.1
## [34,] 158997.12 171176.8 183356.5
## [35,] 156927.45 169053.5 181179.5
## [36,] 154794.08 166867.6 178941.2
## [37,] 152605.00 164627.9 176650.8
## [38,] 150368.67 162343.6 174318.5
## [39,] 148095.28 160025.6 171956.0
## [40,] 145795.70 157685.8 169575.8
## [41,] 143481.49 155336.5 167191.5
## [42,] 141164.76 152990.7 164816.6
## [43,] 138857.85 150661.6 162465.4
## [44,] 136573.14 148362.4 160151.7
## [45,] 134322.76 146105.9 157889.1
## [46,] 132118.32 143904.4 155690.5
## [47,] 129970.65 141769.1 153567.6
## [48,] 127889.62 139710.3 151531.1
## [49,] 125883.88 137736.8 149589.8
## [50,] 123960.66 135855.8 147751.0
## [51,] 122125.72 134072.9 146020.1
## [52,] 120383.59 132392.3 144401.0
## [53,] 118737.40 130816.7 142896.1
## [54,] 117188.65 129347.0 141505.4
## [55,] 115737.57 127982.7 140227.7
## [56,] 114383.22 126721.9 139060.5
## [57,] 113123.58 125561.8 138000.0
## [58,] 111955.75 124498.6 137041.4
## [59,] 110876.09 123527.7 136179.3
## [60,] 109880.40 122644.0 135407.7
## [61,] 108964.11 121842.2 134720.2
## [62,] 108122.40 121116.4 134110.5
## [63,] 107350.29 120461.2 133572.1

```

```
## [64,] 106642.06 119870.2 133098.3
## [65,] 105996.08 119341.1 132686.2
## [66,] 105406.42 118867.7 132328.9
## [67,] 104866.59 118442.3 132018.0
## [68,] 104373.47 118061.9 131750.3
## [69,] 103922.08 117721.3 131520.5
## [70,] 103511.49 117419.3 131327.1
## [71,] 103139.05 117152.6 131166.1
## [72,] 102796.32 116912.4 131028.5
## [73,] 102482.09 116697.6 130913.1
## [74,] 102197.31 116509.3 130821.2
## [75,] 101934.73 116339.9 130745.0
## [76,] 101693.90 116187.9 130681.8
## [77,] 101472.07 116051.9 130631.6
## [78,] 101265.93 115927.2 130588.5
## [79,] 101074.12 115813.4 130552.8
## [80,] 100895.41 115708.4 130521.5
## [81,] 100727.89 115610.7 130493.5
## [82,] 100567.47 115516.2 130464.9
## [83,] 100417.38 115427.5 130437.6
## [84,] 100272.99 115340.7 130408.5
## [85,] 100132.01 115253.6 130375.1
## [86,] 99995.72 115167.1 130338.4
## [87,] 99864.31 115081.5 130298.8
## [88,] 99734.43 114994.0 130253.6
## [89,] 99605.68 114904.3 130202.9
## [90,] 99478.24 114812.8 130147.3
## [91,] 99352.11 114719.7 130087.3
## [92,] 99227.07 114625.1 130023.0
## [93,] 99100.68 114527.2 129953.6
## [94,] 98976.99 114429.7 129882.4
## [95,] 98851.93 114329.8 129807.7
## [96,] 98728.10 114230.0 129731.9
## [97,] 98605.97 114131.2 129656.4
## [98,] 98484.38 114032.8 129581.3
## [99,] 98364.96 113936.9 129508.8
## [100,] 98285.04 113880.2 129475.3
```

```
plot(cv.r)
```



```
log(cv.r$lambda.min)
```

```
## [1] 3.239784
```

```
log(cv.r$lambda.1se)
```

```
## [1] 7.519336
```

```
which.min(cv.r$lambda)
```

```
## [1] 100
```

```
which(cv.r$lambda==cv.r$lambda.1se)
```

```
## [1] 54
```

```
b.min <- predict(cv.r, type="coefficients", s=cv.r$lambda.min)
b.1se <- predict(cv.r, type="coefficients", s=cv.r$lambda.1se)
cbind(b.min, b.1se)
```

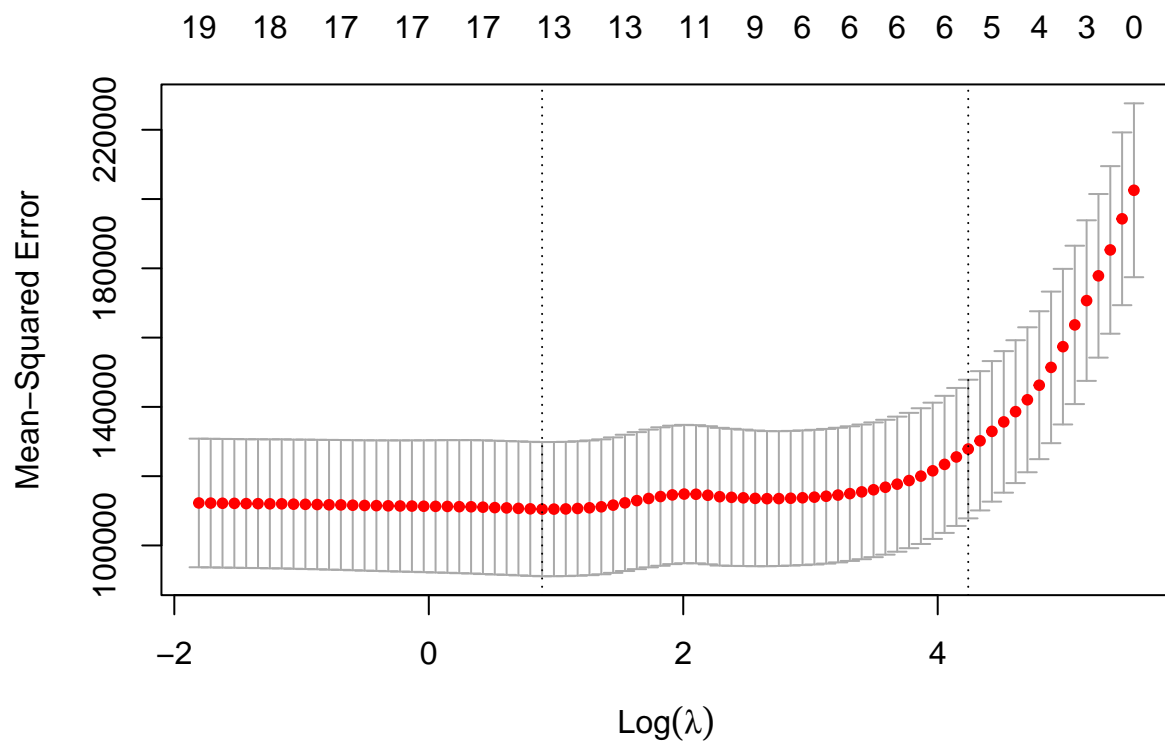
```
## 20 x 2 sparse Matrix of class "dgCMatrix"
##           s1      s1
```

```
## (Intercept)    81.126932 159.796625
## AtBat         -100.212924  15.067887
## Hits          124.863402  20.125453
## HmRun         -11.936650  11.266961
## Runs          25.869160  17.918165
## RBI           18.419837  17.744128
## Walks         73.236073  20.071835
## Years        -43.380051  12.988099
## CAtBat        -2.737486  19.961033
## CHits         88.053957  22.229483
## CHmRun        57.264389  20.805218
## CRuns         97.811727  22.767565
## CRBI          82.970281  23.023410
## CWalks       -73.522590  17.424817
## LeagueN       26.563051   2.693852
## DivisionW    -61.406126 -14.545703
## PutOuts       73.730681  18.945085
## Assists       24.599326   1.332491
## Errors       -24.303154  -1.556113
## NewLeagueN   -9.028894   2.222950
```

```
c(sqrt(sum(b.min[-1]^2)), sqrt(sum(b.1se[-1]^2)))
```

```
## [1] 279.4935  72.3250
```

```
set.seed(2)
cv.l <- cv.glmnet(x, y, alpha=1, nfolds=10)
plot(cv.l)
```



```
log(cv.l$lambda.min)
```

```
## [1] 0.8906821
```

```
log(cv.l$lambda.1se)
```

```
## [1] 4.239897
```

```
which(cv.l$lambda==cv.l$lambda.min)
```

```
## [1] 51
```

```
which(cv.l$lambda==cv.l$lambda.1se)
```

```
## [1] 15
```

```
b.min <- predict(cv.l, type="coefficients", s=cv.l$lambda.min)
b.1se <- predict(cv.l, type="coefficients", s=cv.l$lambda.1se)
cbind(b.min, b.1se)
```

```
## 20 x 2 sparse Matrix of class "dgCMatrix"
##           s1           s1
```

```
## (Intercept) 129.41556 127.956948
## AtBat      -237.15665 .
## Hits       261.49419 64.110315
## HmRun      . .
## Runs       . .
## RBI        . .
## Walks      105.06567 34.295642
## Years      -47.71291 .
## CAtBat     . .
## CHits      . .
## CHmRun     44.09343 .
## CRuns      225.18108 52.983406
## CRBI       125.98840 108.663331
## CWalks     -146.53946 .
## LeagueN    16.20588 .
## DivisionW  -59.66326 -4.030128
## PutOuts    76.60907 23.451888
## Assists    26.87539 .
## Errors     -14.27657 .
## NewLeagueN . .
```

```
c(sum(abs(b.min[-1])), sum(abs(b.1se[-1])))
```

```
## [1] 1386.8620 287.5347
```

Ridge: The Bias-Variance tradeoff

```
set.seed(1234)
K <- 100
p <- 40
n <- 50
beta <- runif(p, -1, 1)
lam <- 10^seq(3, -3, length.out=50)
x <- matrix(rnorm(n * p), n, p)
bhat <- array(0, c(p, length(lam), K))

for (i in 1:K) {
  y <- x %*% beta + rnorm(n)
  fit <- glmnet(x, y, alpha=0, lambda=lam)
  bhat[,i] <- as.matrix(fit$beta)
}

MSE0 <- Bias0 <- Vars0 <- matrix(0, p, length(lam))

for (k in 1:length(lam)) {
  MS <- (bhat[,k] - matrix(beta, p, K))^2
  MSE0[,k] <- apply(MS, 1, mean)
  Bias0[,k] <- abs(apply(bhat[,k], 1, mean) - beta)
  Vars0[,k] <- apply(bhat[,k], 1, var)
}
```

```
MSE <- apply(MSE0, 2, mean)
Bias <- apply(Bias0, 2, mean)
Vars <- apply(Vars0, 2, mean)
```

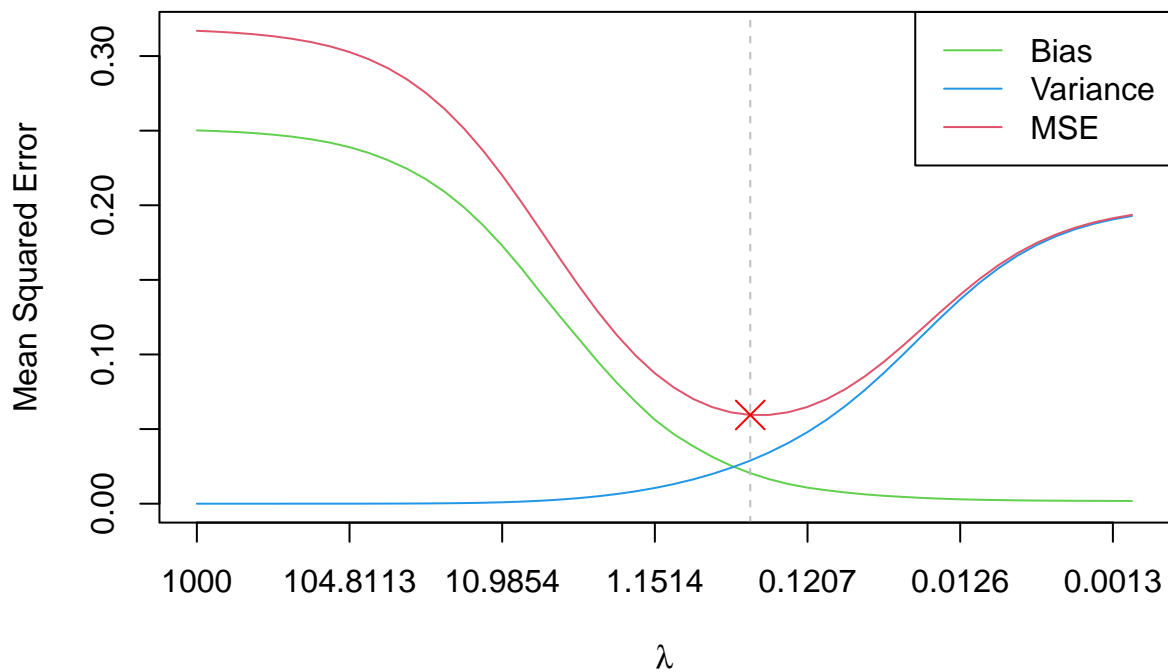
```
MAT <- cbind(Bias^2, Vars, MSE)
data.frame(lambda=round(lam, 3), MSE=MAT[,3])
```

##	lambda	MSE
## 1	1000.000	0.31693101
## 2	754.312	0.31635426
## 3	568.987	0.31559416
## 4	429.193	0.31459428
## 5	323.746	0.31328221
## 6	244.205	0.31156599
## 7	184.207	0.30933045
## 8	138.950	0.30643403
## 9	104.811	0.30270702
## 10	79.060	0.29795273
## 11	59.636	0.29195354
## 12	44.984	0.28448407
## 13	33.932	0.27533318
## 14	25.595	0.26433534
## 15	19.307	0.25140859
## 16	14.563	0.23659337
## 17	10.985	0.22008324
## 18	8.286	0.20223746
## 19	6.251	0.18356741
## 20	4.715	0.16469051
## 21	3.556	0.14626235
## 22	2.683	0.12889911
## 23	2.024	0.11309967
## 24	1.526	0.09920968
## 25	1.151	0.08742132
## 26	0.869	0.07778465
## 27	0.655	0.07025636
## 28	0.494	0.06475055
## 29	0.373	0.06117498
## 30	0.281	0.05945335
## 31	0.212	0.05950986
## 32	0.160	0.06131919
## 33	0.121	0.06484610
## 34	0.091	0.07005856
## 35	0.069	0.07691802
## 36	0.052	0.08531573
## 37	0.039	0.09508516
## 38	0.029	0.10590440
## 39	0.022	0.11733858
## 40	0.017	0.12888429
## 41	0.013	0.14005615
## 42	0.010	0.15047956
## 43	0.007	0.15980298
## 44	0.005	0.16789986
## 45	0.004	0.17472040


```
## 46    0.003 0.18032160
## 47    0.002 0.18489332
## 48    0.002 0.18848420
## 49    0.001 0.19134091
## 50    0.001 0.19354913
```

```
matplot(MAT, type="l", col=c(3,4,2), lty=1, xaxt="n",
        xlab=expression(lambda), ylab="Mean Squared Error")
legend("topright", c("Bias", "Variance", "MSE"), col=c(3,4,2),
       lty=1)
```

```
w <- which.min(MAT[,3])
abline(v=w, col="gray", lty=2)
points(w, MAT[w, 3], pch=4, col="red", cex=2)
cc <- seq(1, 50, 8)
axis(1, at=cc, labels=round(lam[cc],4))
```



Lasso: The Bias-Variance Trade Off

```
set.seed(1234)
K <- 100
p <- 40
n <- 50
```

```

beta <- runif(p, -1, 1)
lam <- 10^seq(0.7, -4, length.out=50)
x <- matrix(rnorm(n * p), n, p)
bhat <- array(0, c(p, length(lam), K))
for (i in 1:K) {
  y <- x %*% beta + rnorm(n)
  fit <- glmnet(x, y, alpha=1, lambda=lam)
  bhat[,i] <- as.matrix(fit$beta)
}
MSE0 <- Bias0 <- Vars0 <- matrix(0, p, length(lam))
for (k in 1:length(lam)) {
  MS <- (bhat[,k,] - matrix(beta, p, K))^2
  MSE0[,k] <- apply(MS, 1, mean)
  Bias0[,k] <- abs(apply(bhat[,k,], 1, mean) - beta)
  Vars0[,k] <- apply(bhat[,k,], 1, var)
}

```

```

MSE <- apply(MSE0, 2, mean)
Bias <- apply(Bias0, 2, mean)
Vars <- apply(Vars0, 2, mean)
MAT <- cbind(Bias^2, Vars, MSE)
data.frame(lambda=round(lam, 3), MSE=MAT[,3])

```

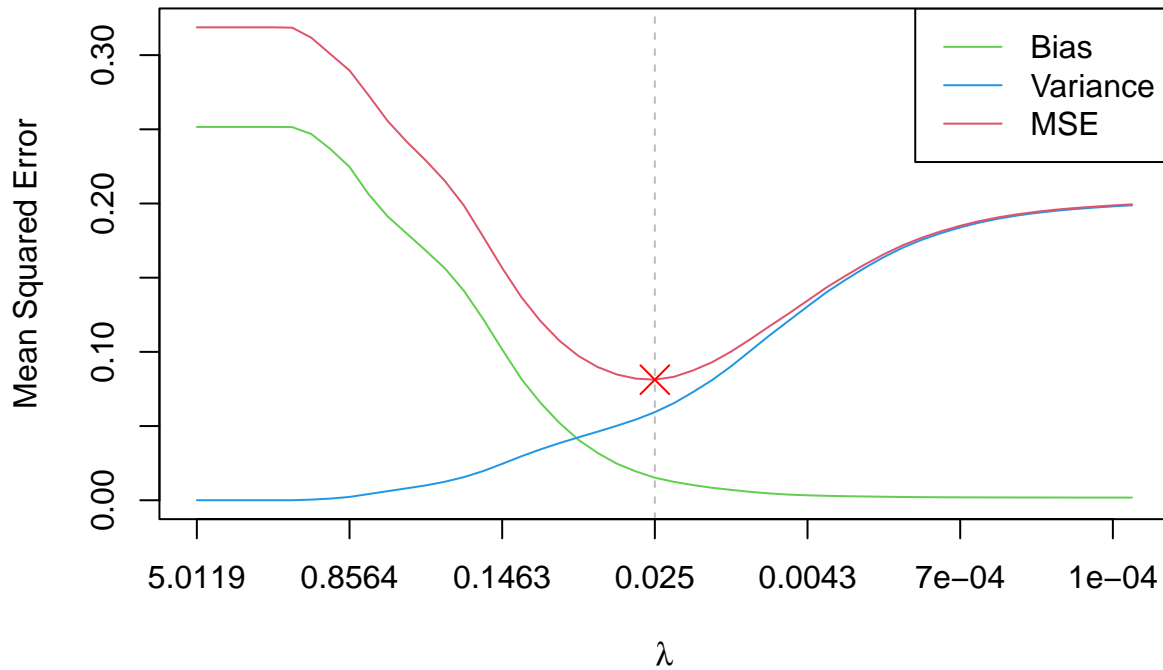
```

##      lambda      MSE
## 1  5.012 0.31872022
## 2  4.019 0.31872022
## 3  3.222 0.31872022
## 4  2.584 0.31872022
## 5  2.072 0.31872022
## 6  1.661 0.31850322
## 7  1.332 0.31171991
## 8  1.068 0.30066269
## 9  0.856 0.28965552
## 10 0.687 0.27292256
## 11 0.551 0.25560874
## 12 0.441 0.24162140
## 13 0.354 0.22897777
## 14 0.284 0.21512478
## 15 0.228 0.19851600
## 16 0.182 0.17771447
## 17 0.146 0.15656136
## 18 0.117 0.13699184
## 19 0.094 0.12084441
## 20 0.075 0.10749690
## 21 0.060 0.09720984
## 22 0.048 0.08979449
## 23 0.039 0.08479431
## 24 0.031 0.08197614
## 25 0.025 0.08125365
## 26 0.020 0.08319555
## 27 0.016 0.08748459
## 28 0.013 0.09301937
## 29 0.010 0.10024209

```

```
## 30 0.008 0.10850368
## 31 0.007 0.11722004
## 32 0.005 0.12570653
## 33 0.004 0.13455070
## 34 0.003 0.14345889
## 35 0.003 0.15135923
## 36 0.002 0.15890115
## 37 0.002 0.16579777
## 38 0.001 0.17190427
## 39 0.001 0.17699509
## 40 0.001 0.18128398
## 41 0.001 0.18497898
## 42 0.001 0.18810408
## 43 0.000 0.19067300
## 44 0.000 0.19273950
## 45 0.000 0.19442198
## 46 0.000 0.19582081
## 47 0.000 0.19694150
## 48 0.000 0.19786052
## 49 0.000 0.19871040
## 50 0.000 0.19941652
```

```
matplot(MAT, type="l", col=c(3,4,2), lty=1, xaxt="n",
xlab=expression(lambda), ylab="Mean Squared Error")
legend("topright", c("Bias", "Variance", "MSE"), col=c(3,4,2),
lty=1)
w <- which.min(MAT[,3])
abline(v=w, col="gray", lty=2)
points(w, MAT[w, 3], pch=4, col="red", cex=2)
cc <- seq(1, 50, 8)
axis(1, at=cc, labels=round(lam[cc],4))
```



```

MSE.fun <- function(n, p, K, beta, lam, xtest, ytest) {
  yhat0 <- yhat1 <- array(0, c(n, length(lam), K))
  for (i in 1:K) {
    x <- matrix(rnorm(n * p), n, p)
    y <- x %*% beta + rnorm(n)
    g0 <- glmnet(x, y, alpha=0, lambda=lam)
    g1 <- glmnet(x, y, alpha=1, lambda=lam)
    yhat0[1:n, 1:length(lam), i] <- predict(g0, x.test)
    yhat1[1:n, 1:length(lam), i] <- predict(g1, x.test)
  }
  MSE0 <- Bias0 <- Vars0 <- array(0, c(n, length(lam)))
  MSE1 <- Bias1 <- Vars1 <- array(0, c(n, length(lam)))
  for (j in 1:length(lam)) {
    PE0 <- (yhat0[,j,] - matrix(ytest, n, K))^2
    PE1 <- (yhat1[,j,] - matrix(ytest, n, K))^2
    MSE0[,j] <- apply(PE0, 1, mean)
    MSE1[,j] <- apply(PE1, 1, mean)
    BS0 <- abs(yhat0[,j,] - matrix(ytest, n, K))
    BS1 <- abs(yhat1[,j,] - matrix(ytest, n, K))
    Bias0[,j] <- apply(BS0, 1, mean)
    Bias1[,j] <- apply(BS1, 1, mean)
    Vars0[,j] <- apply(yhat0[,j,], 1, var)
    Vars1[,j] <- apply(yhat1[,j,], 1, var)
  }
  MSE.r <- apply(MSE0, 2, mean)
  MSE.l <- apply(MSE1, 2, mean)

```

```

Bia.r <- apply(Bias0, 2, mean)
Bia.l <- apply(Bias1, 2, mean)
Var.r <- apply(Vars0, 2, mean)
Var.l <- apply(Vars1, 2, mean)
ridge <- apply(cbind(Bia.r^2, Var.r, MSE.r), 2, rev)
lasso <- apply(cbind(Bia.l^2, Var.l, MSE.l), 2, rev)
newlam <- rev(lam)
return(list(ridge=ridge,lasso=lasso, lambda=newlam))
}

```

```

set.seed(111000)
K <- 10
p <- 120
n <- 100
lam <- 10^seq(1, -3, -0.05)
x.test <- matrix(rnorm(n * p), n, p)

```

The case that all predictors have non-zero coefficients

```

beta1 <- beta2 <- runif(p, -1, 1)
ytest1 <- x.test %*% beta1 + rnorm(n)
g1 <- MSE.fun(n, p, K, beta1, lam, xtest, ytest1)
RES1 <- cbind(g1$lasso, g1$ridge)

```

The case that only 5 predictors have non-zero coefficients

```

beta2[6:p] <- 0
ytest2 <- x.test %*% beta2 + rnorm(n)
g2 <- MSE.fun(n, p, K, beta2, lam, xtest, ytest2)
RES2 <- cbind(g2$lasso, g2$ridge)

```

```

par(mfrow=c(1,2))
matplot(RES1, type="l", col=c(1,3,2), lty=rep(1:2,each=3),
xlab=expression(lambda), xaxt="n",
ylab="Mean Squared Error")
cc <- c(1, seq(21, 81, 20))
axis(1, at=cc, labels=g1$lambda[cc])
matplot(RES2, type="l", col=c(1,3,2), lty=rep(1:2,each=3),
xlab=expression(lambda), xaxt="n",
ylab="Mean Squared Error")
legend("topright",c("Bias_Lasso", "Variance_Lasso", "MSE_Lasso",
"Bias_Ridge", "Variance_Ridge", "MSE_Ridge"),
col=c(1,3,2), lty=rep(1:2, each=3))
axis(1, at=cc, labels=g2$lambda[cc])

```

